EXPLAINABLE MACHINE LEARNING

Hung-yi Lee 李宏毅

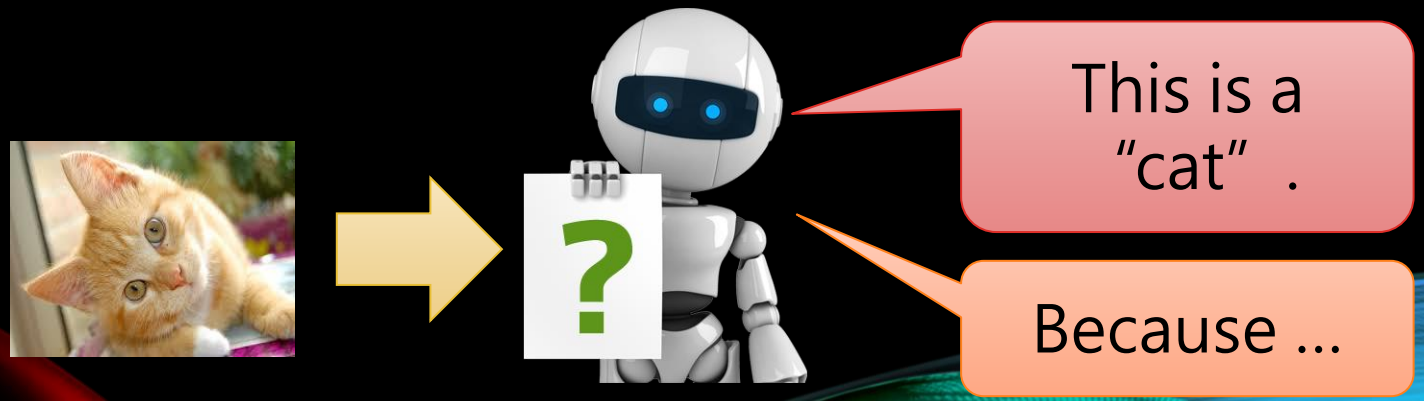

Why we need Explainable ML?

- Correct answers ≠ Intelligent

有一隻馬
他會在眾人面前算數學
例如簡易的加法
（他會用踩踏次數給出答案）
但他其實不是真的會算
而是看周遭人們的反應
而去看什麼時候要停止踩踏

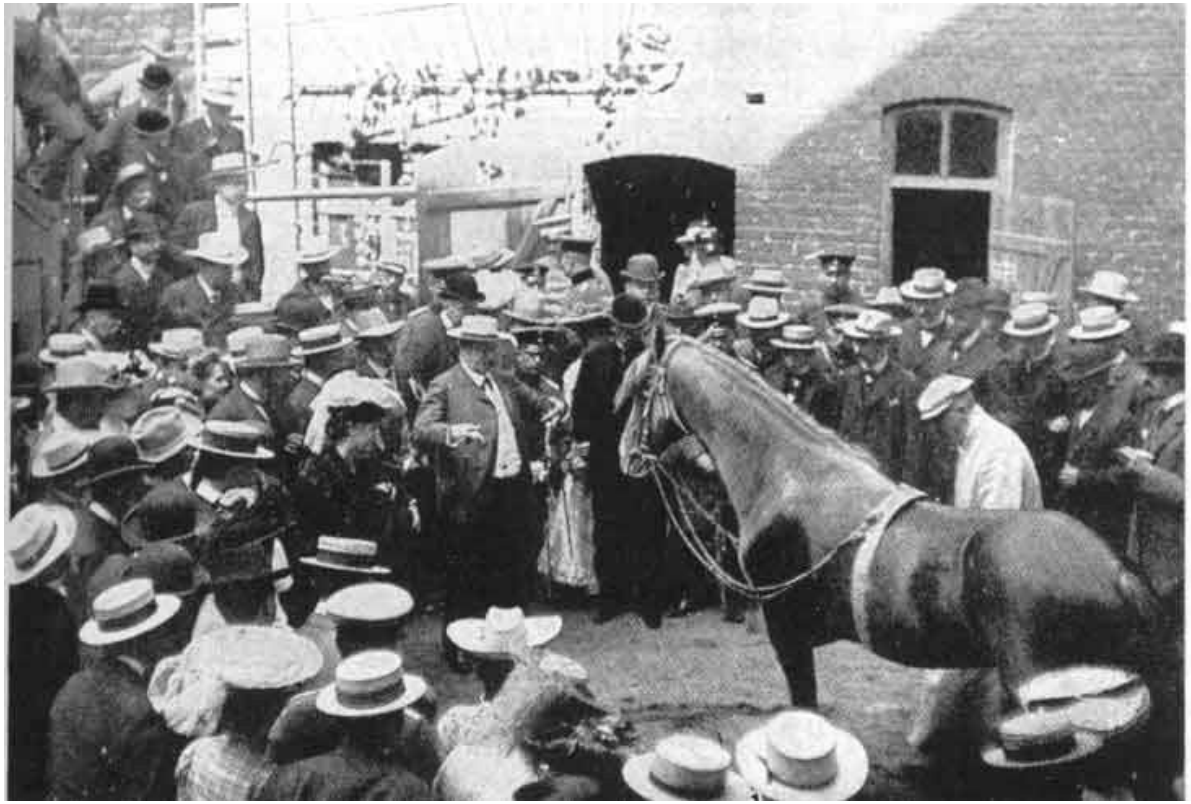

Why we need Explainable ML?

- Loan issuers are required by law to explain their models.
- Medical diagnosis model is responsible for human life. Can it be a black box?
- If a model is used at the court, we must make sure the model behaves in a nondiscriminatory manner.
- If a self-driving car suddenly acts abnormally, we need to explain why.

We can improve ML model based on explanation.

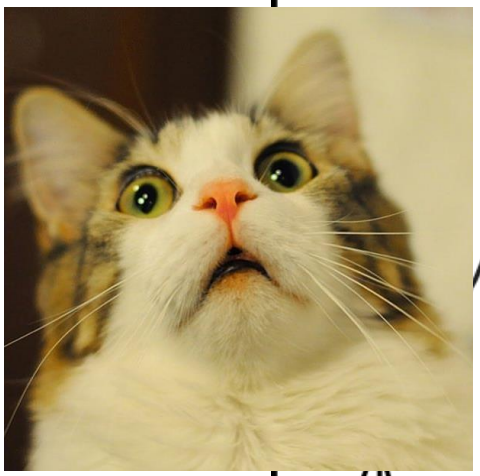

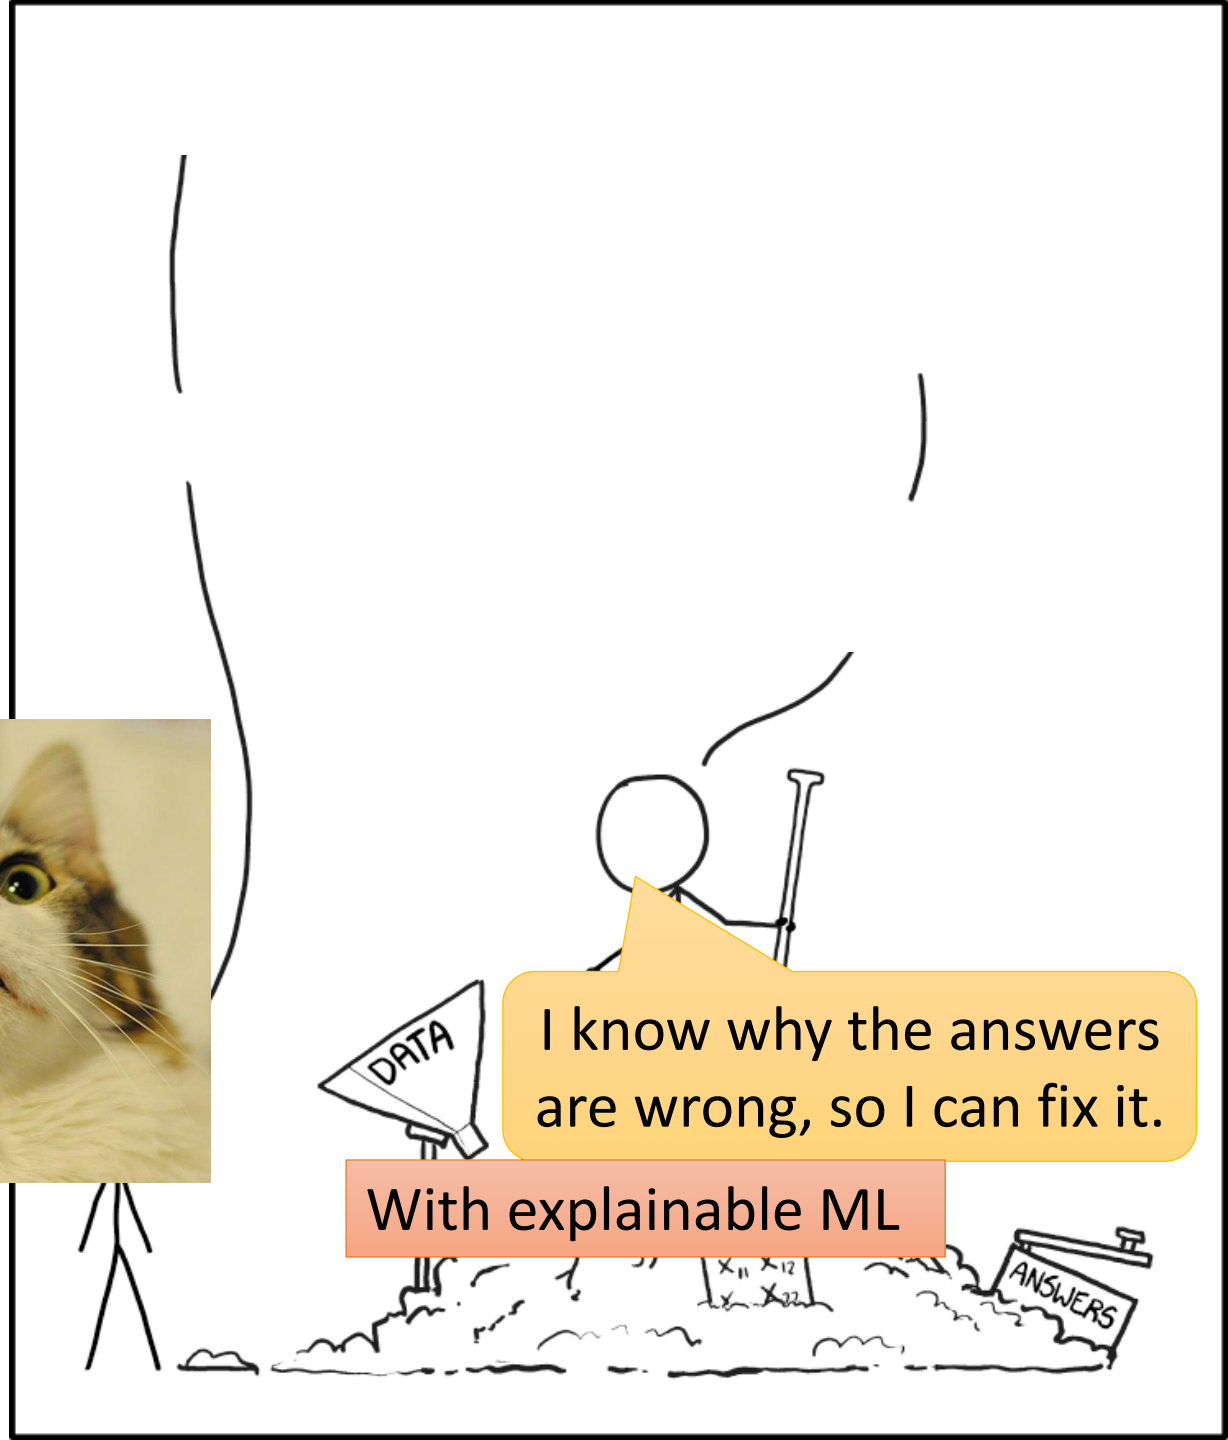

https://www.explainxkcd.com/wiki/index.php/1838:_Machine_Learning

Interpretable v.s. Powerful

- Some models are intrinsically interpretable.
  - For example, linear model (from weights, you know the importance of features)
  - But not very powerful.
- Deep network is difficult to interpretable. Deep networks are black boxes … but powerful than a linear model.

We don't want to use a more powerful model because it is a black box.

This is "cut the feet to fit the shoes." (削足適履)

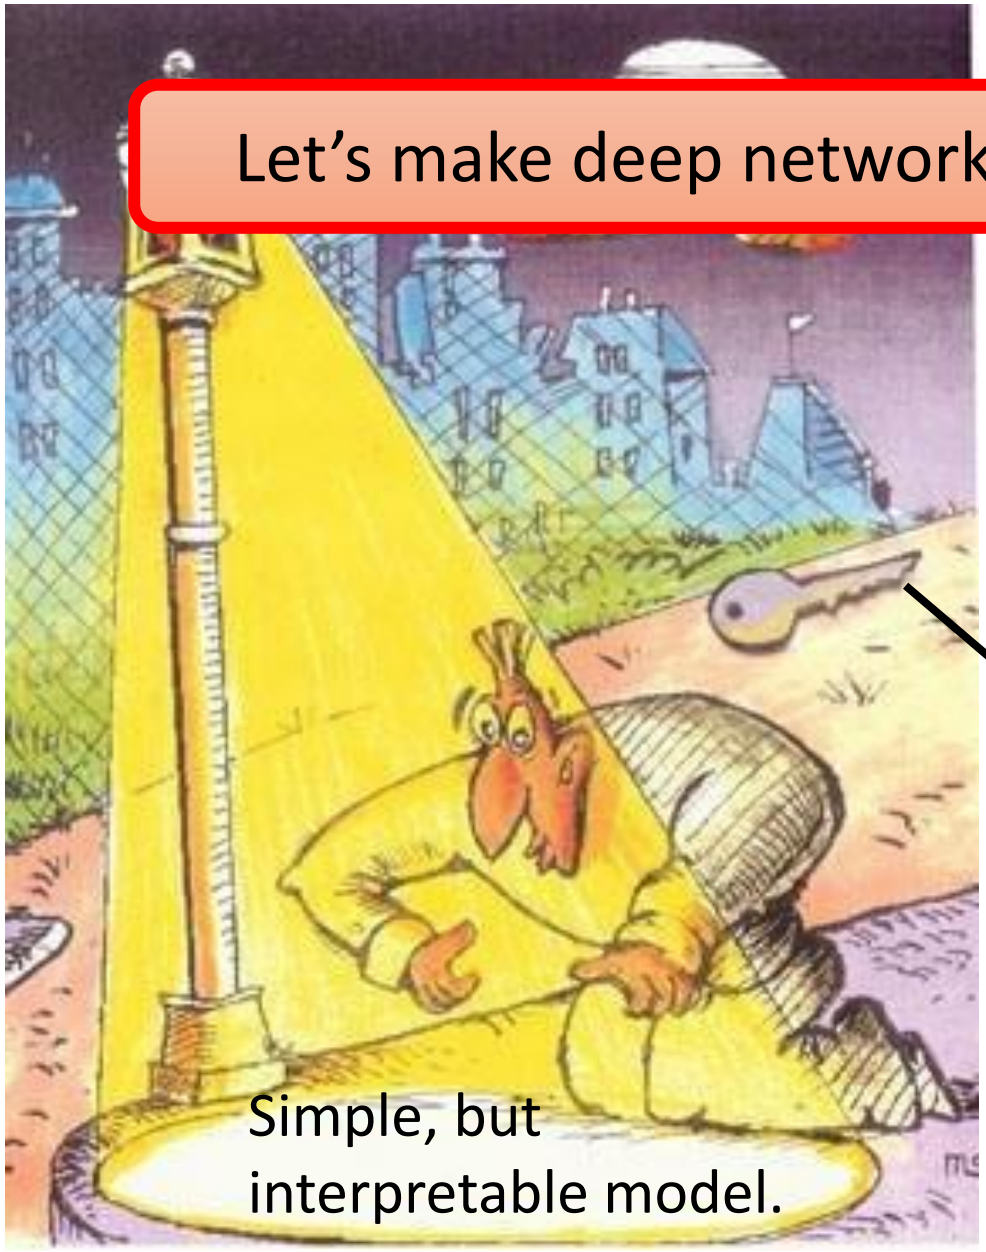

Source of image: https://kknews.cc/news/pnynzgp.html

Interpretable v.s. Powerful

- Are there some models interpretable and powerful at the same time?
- How about decision tree?

Source of image: https://mropengate.blogspot.com/2015/06/ai-ch13-2-decision-tree.html

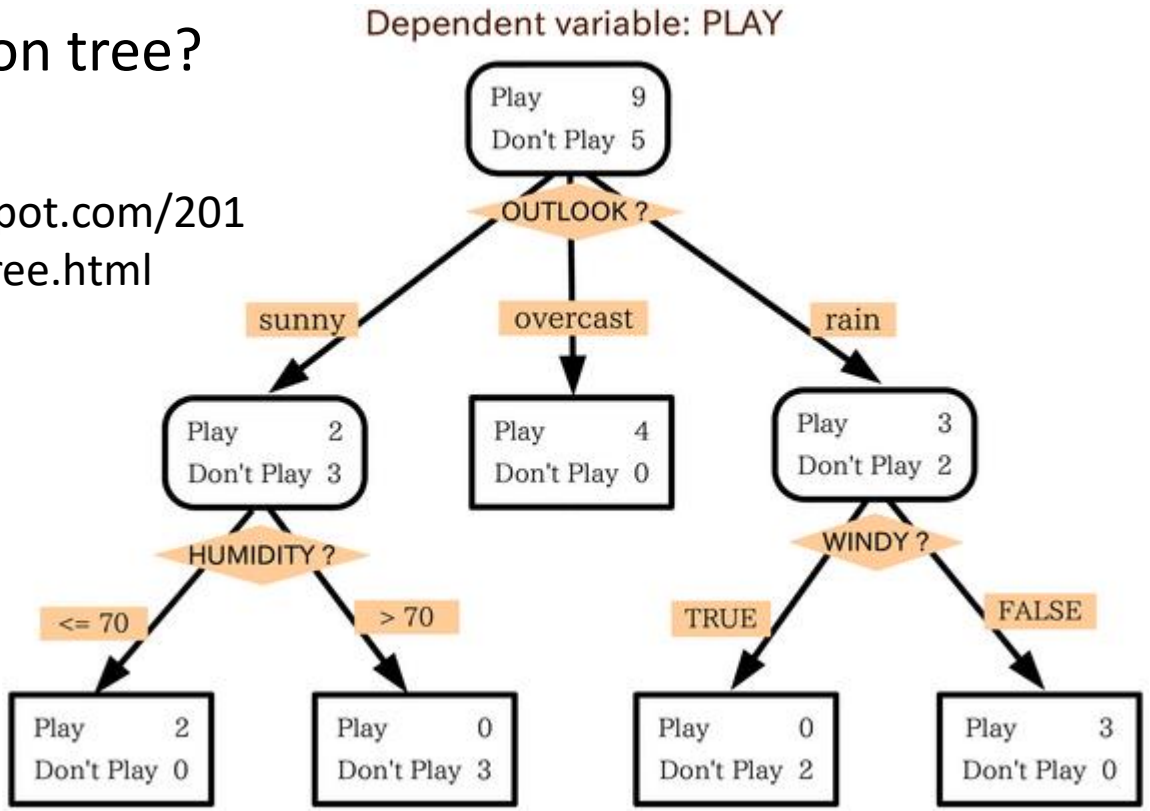


Decision tree is all you need!?

Interpretable v.s. Powerful

- A tree can still be terrible!

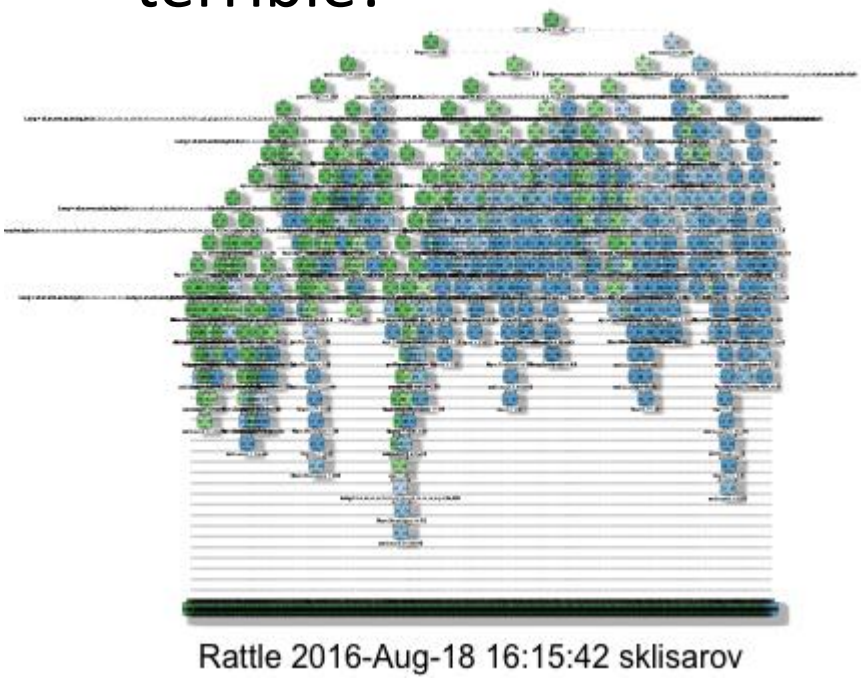

https://stats.stackexchange.com/questions/230581/decision-tree-too-large-to-interpret

- We use a forest!

Goal of Explainable ML

- Completely know how an ML model works?
  - We do not completely know how brains work!
  - But we trust the decision of humans!

The Copy Machine Study (Ellen Langer, Harvard University)

"Excuse me, I have 5 pages. May I use the Xerox machine?"
60% accept

"Excuse me, I have 5 pages. May I use the Xerox machine, **because I'm in a rush**?"
94% accept

"Excuse me, I have 5 pages. May I use the Xerox machine, **because I have to make copies**?"
93% accept

https://jamesclear.com/wp-content/uploads/2015/03/copy-machine-study-ellen-langer.pdf

Make people (your customers, your boss, yourself) comfortable.

(my two cents)

Explainable ML

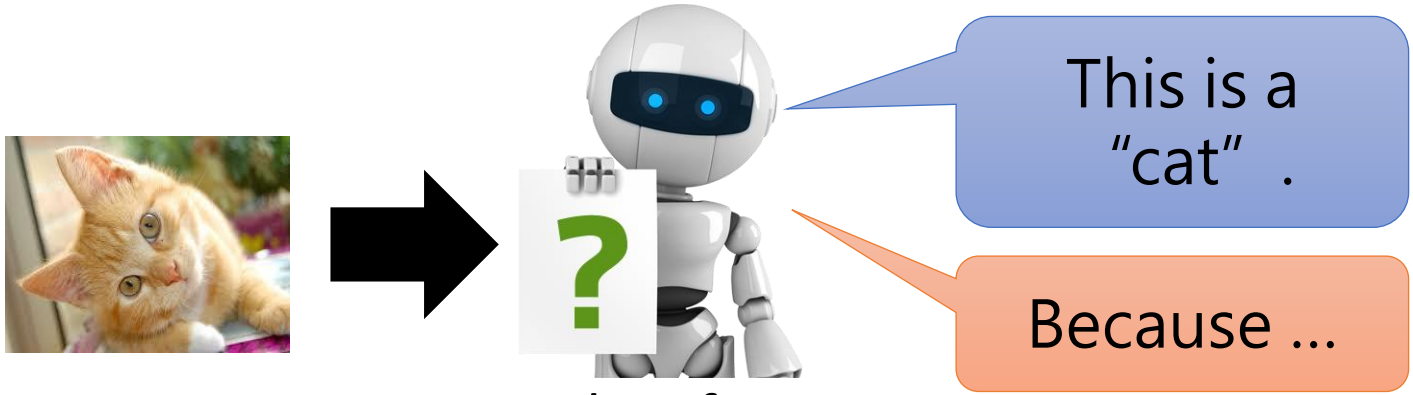

Local Explanation

Why do you think *this image* is a cat?

Global Explanation

What does a "cat" look like?

(not referred to a specific image)

Local Explanation: Explain the Decision

Questions: Why do you think this image is a cat?

Which component is critical?

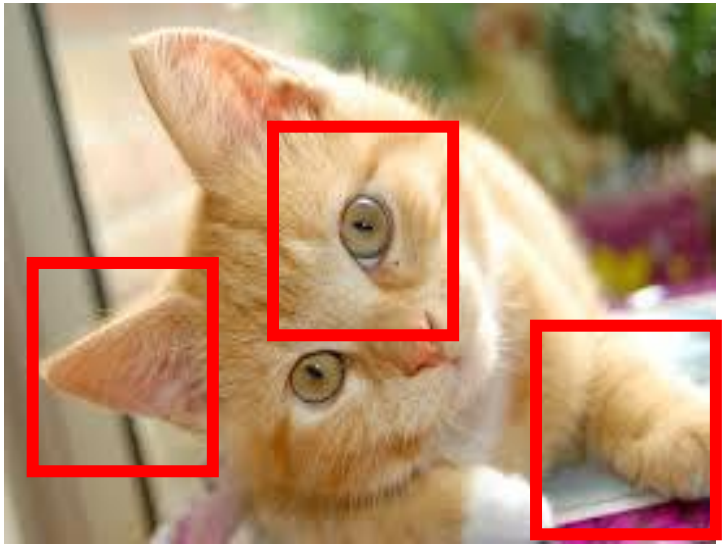

Which component is critical for making decision?

Object x ⟶ Image, text, etc.

Components:

$$\{x_1, \cdots, x_n, \cdots, x_N\}$$

Image: pixel, segment, etc.
Text: a word

- Removing or modifying the components
- Large decision change

➡ Important component

使用灰色框框，然後去覆蓋掉圖片中每一個位置
若model的confidence變低了，就代表機器是看那個地方做決定的

Reference: Zeiler, M. D., & Fergus, R. (2014). Visualizing and understanding convolutional networks. In *Computer Vision–ECCV 2014* (pp. 818-833)

$\{x_1, \cdots, x_n, \cdots, x_N\} \rightarrow \{x_1, \cdots, x_n + \Delta x, \cdots, x_N\}$

pixels

$e \rightarrow e + \Delta e$

$\left|\frac{\Delta e}{\Delta x}\right| \rightarrow \left|\frac{\partial e}{\partial x_n}\right|$

loss of an example (the difference between model output and ground truth)

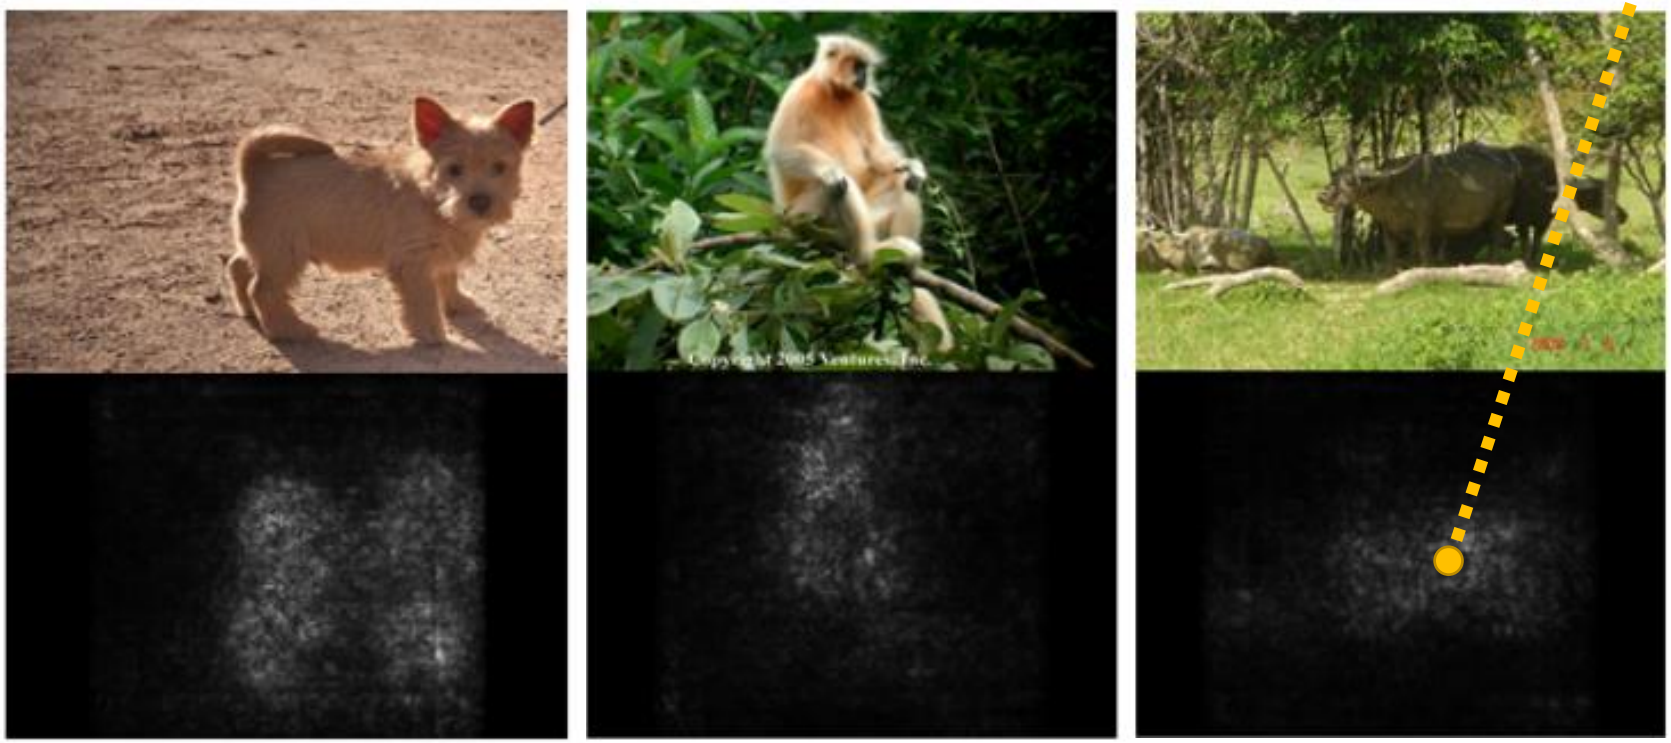

Saliency Map

Karen Simonyan, Andrea Vedaldi, Andrew Zisserman, "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps", ICLR, 2014

Case Study: Pokémon v.s. Digimon

https://medium.com/@tyreeostevenson/teaching-a-computer-to-classify-anime-8c77bc89b881

Task

Pokémon images: https://www.Kaggle.com/kvpratama/pokemon-images-dataset/data

Digimon images: https://github.com/DeathReaper0965/Digimon-Generator-GAN

Pokémon

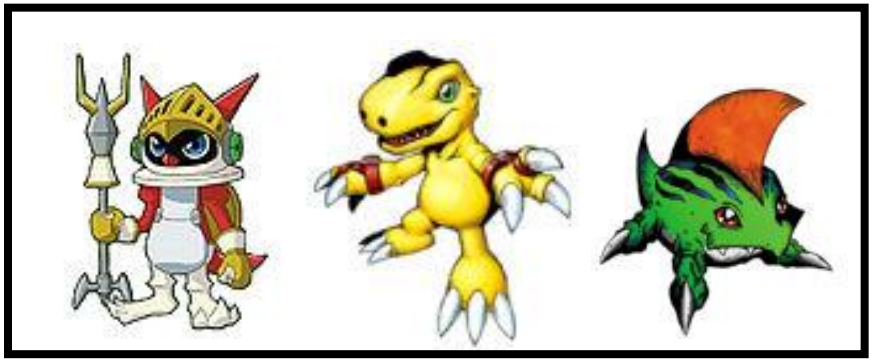

Digimon

Testing Images:

Experimental Results

```
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same', input_shape=(120,120,3)))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(256, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(256, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(1024))
model.add(Activation('relu'))
model.add(Dense(2))
model.add(Activation('softmax'))
```

Training Accuracy: 98.9%

Testing Accuracy: 98.4%

Amazing!!!!!!

Saliency Map

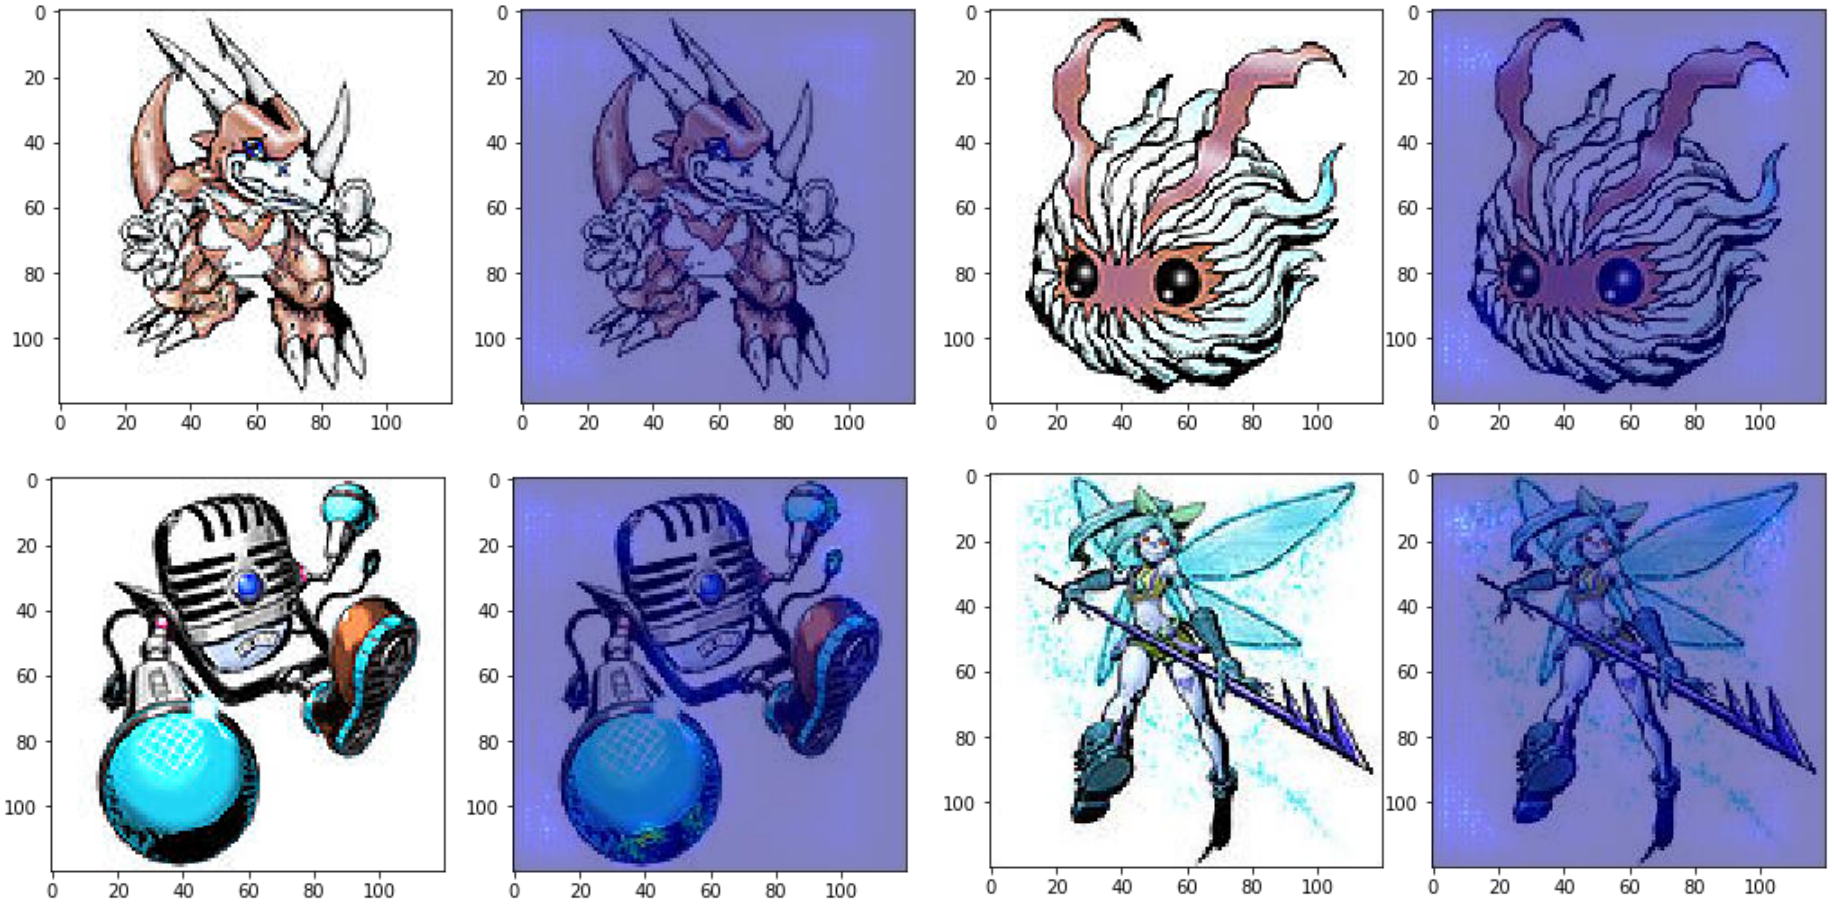

Saliency Map

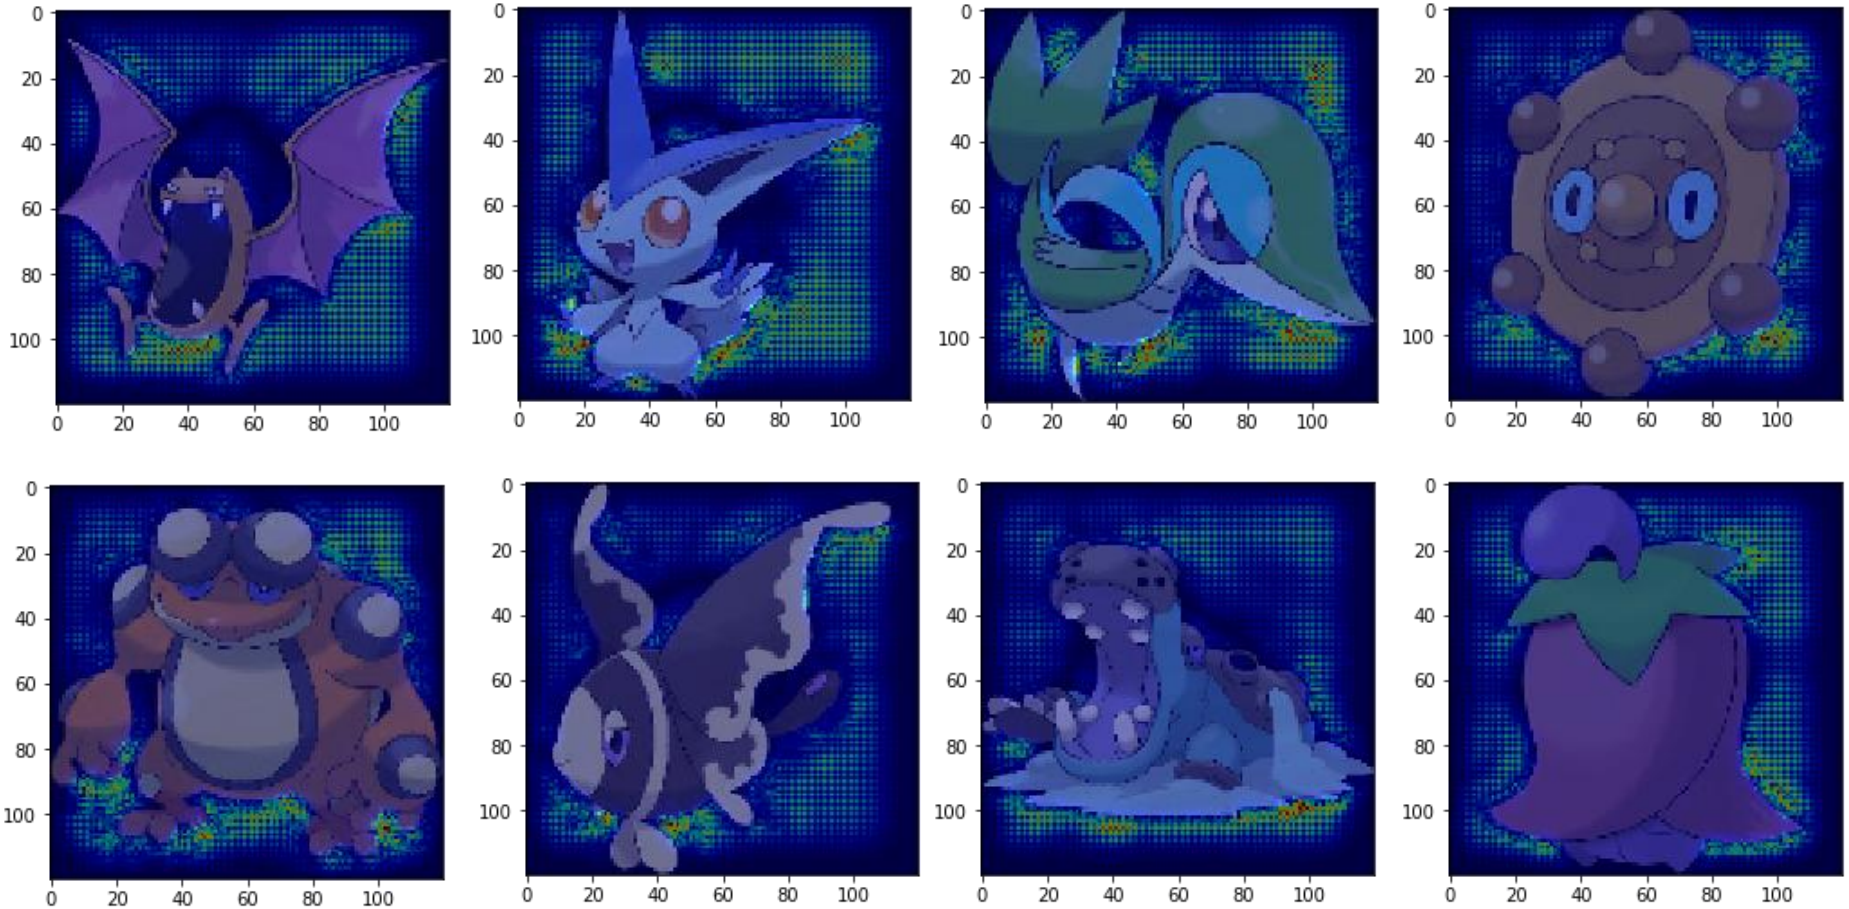

What Happened?

- All the images of Pokémon are PNG, while most images of Digimon are JPEG.

loading the files

png files have transparent background

transparent background becomes black

Machine discriminates Pokémon and Digimon based on the background colors.

More Examples …

- PASCAL VOC 2007 data set

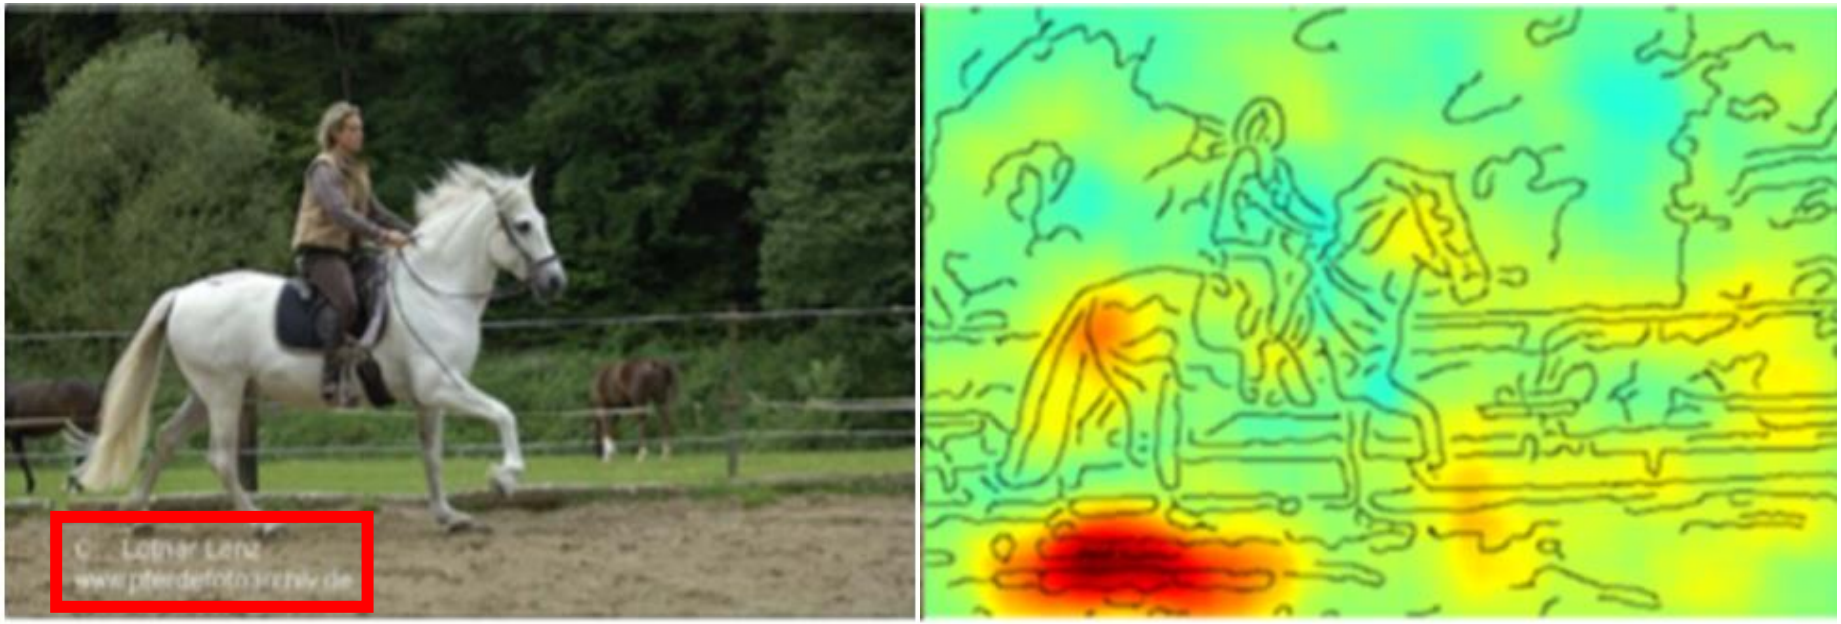

This slide is from: GCPR 2017 Tutorial — W. Samek & K.-R. Müller

Limitation: Noisy Gradient

Gazelle
(瞪羚)

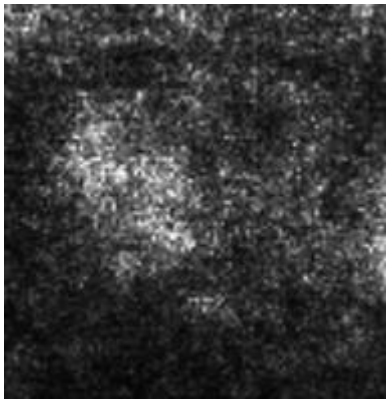

Typical

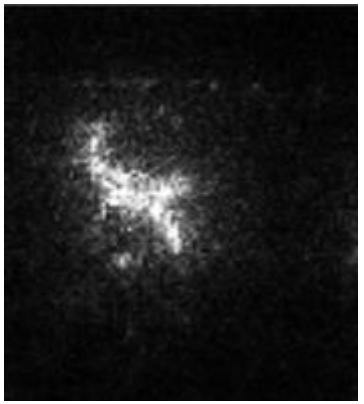

SmoothGrad

SmoothGrad: Randomly add noises to the input image, get saliency maps of the noisy images, and average them.

https://arxiv.org/abs/1706.03825

Limitation: Gradient Saturation

Gradient cannot always reflect importance

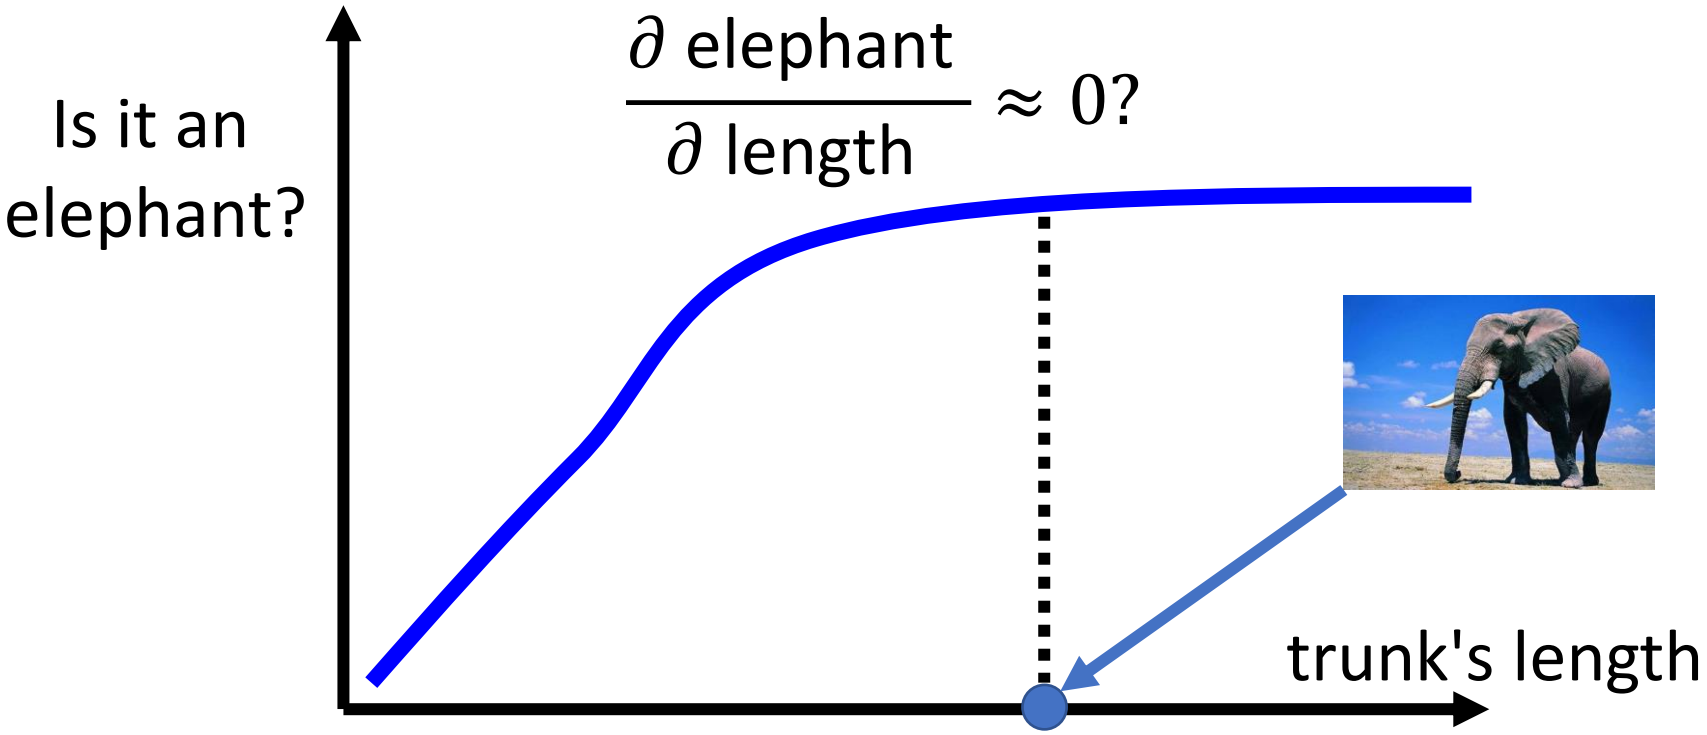

Alternative: Integrated gradient (IG)

https://arxiv.org/abs/1611.02639

How a network processes the input data?

- Visualization

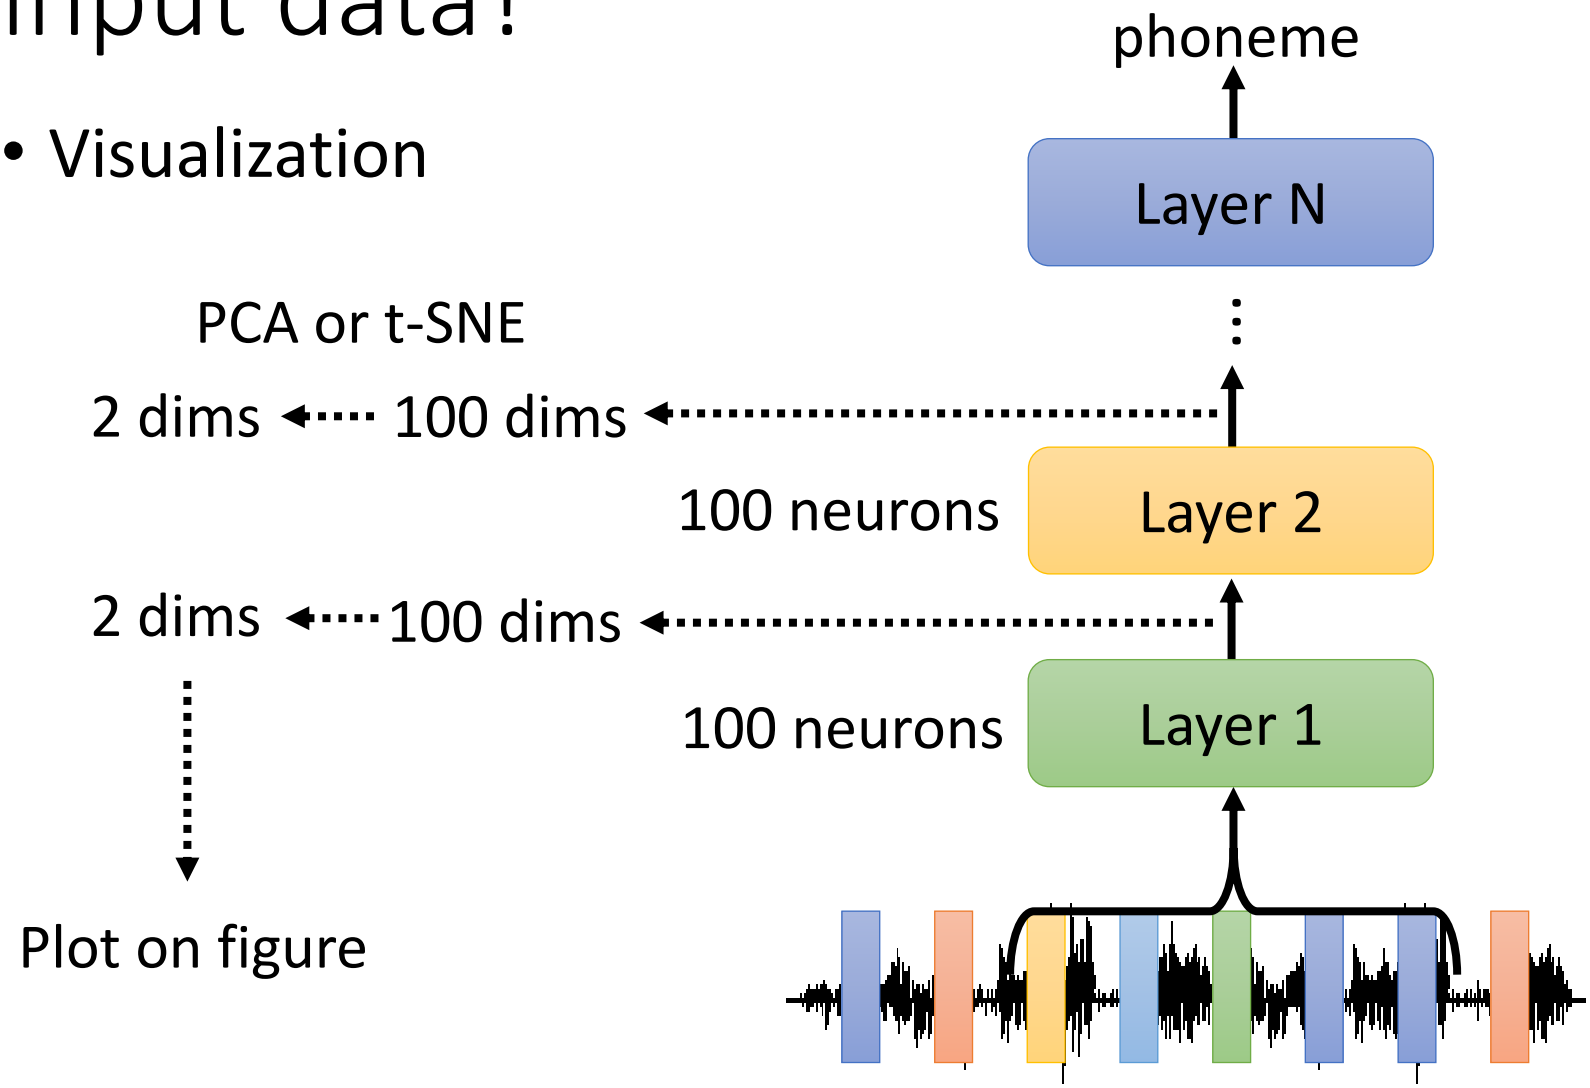

How a network processes the input data?

- Visualization
  Colors: speakers

A. Mohamed, G. Hinton, and G. Penn, "Understanding how Deep Belief Networks Perform Acoustic Modelling," in ICASSP, 2012.

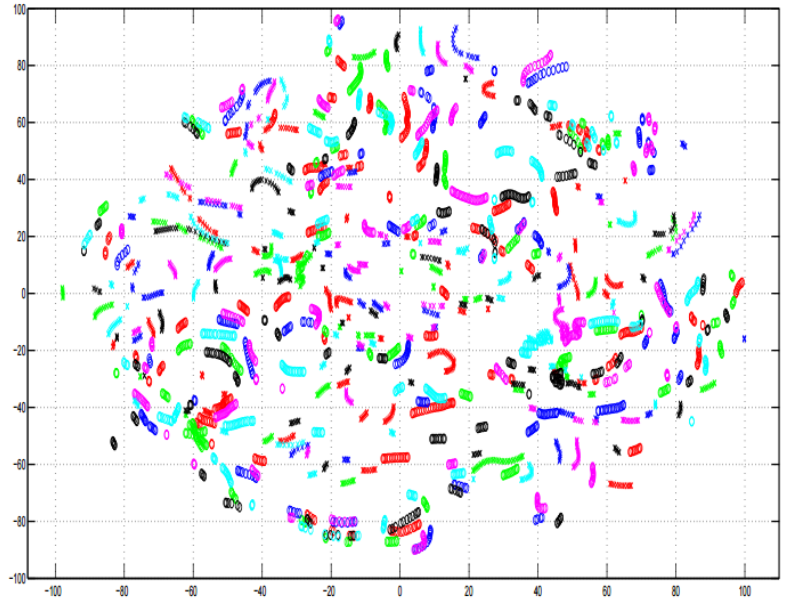

Input Acoustic Feature (MFCC)

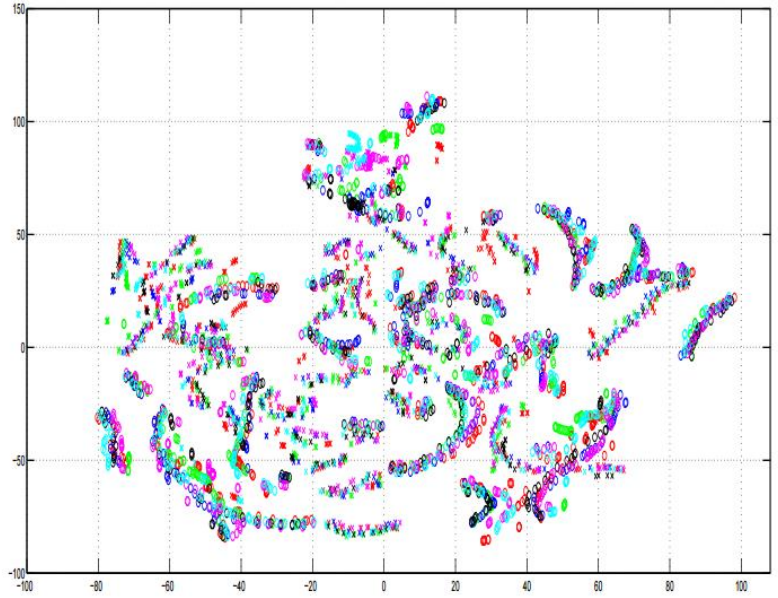

8-th Hidden Layer

How a network processes the input data?

- Visualization

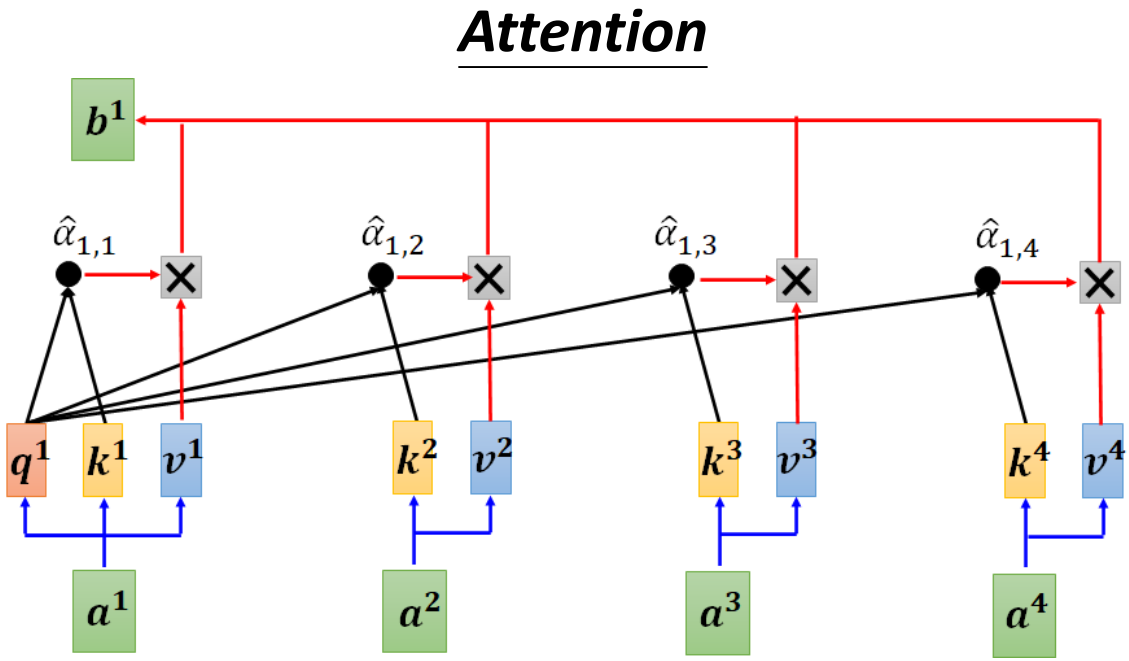

Attention is not Explanation
https://arxiv.org/abs/1902.10186

Attention is not not Explanation
https://arxiv.org/abs/1908.04626

How a network processes the input data?

- Probing

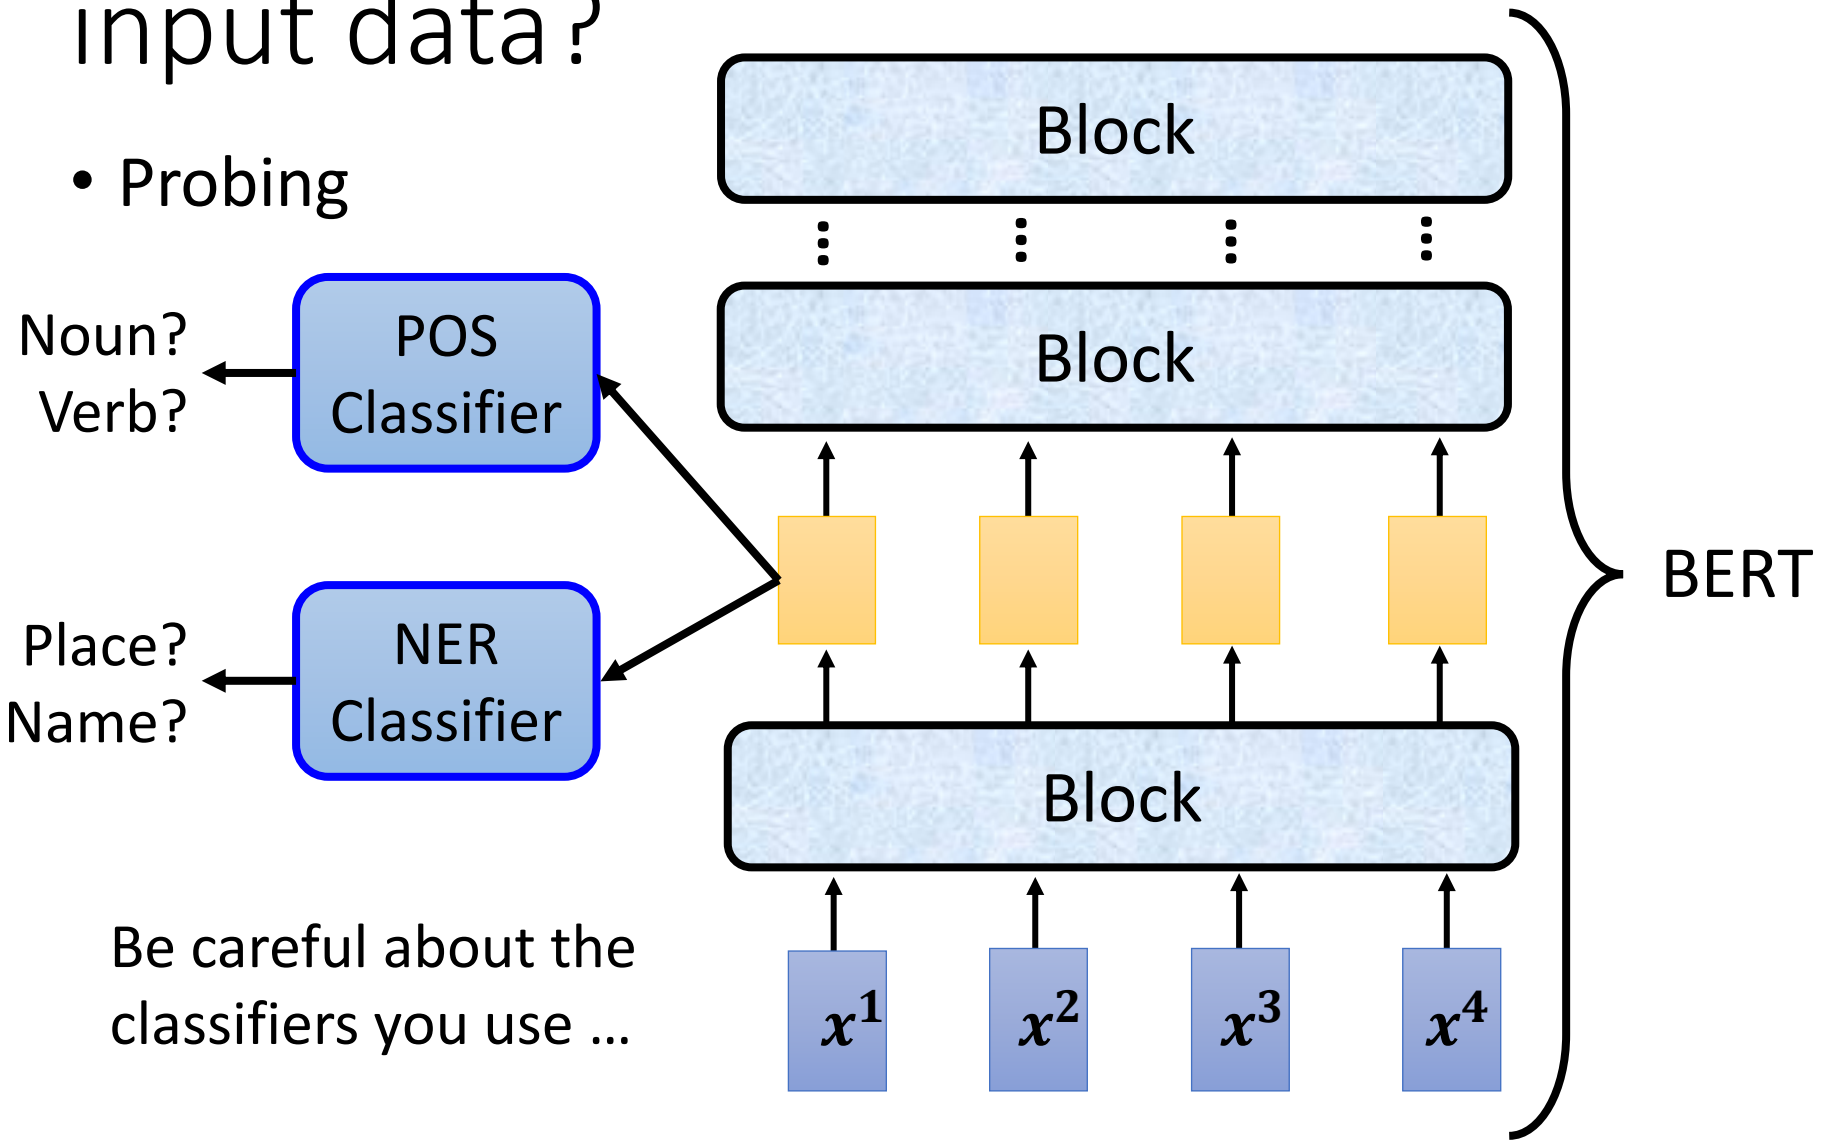

Be careful about the classifiers you use …

How a network processes the input data?

- Probing

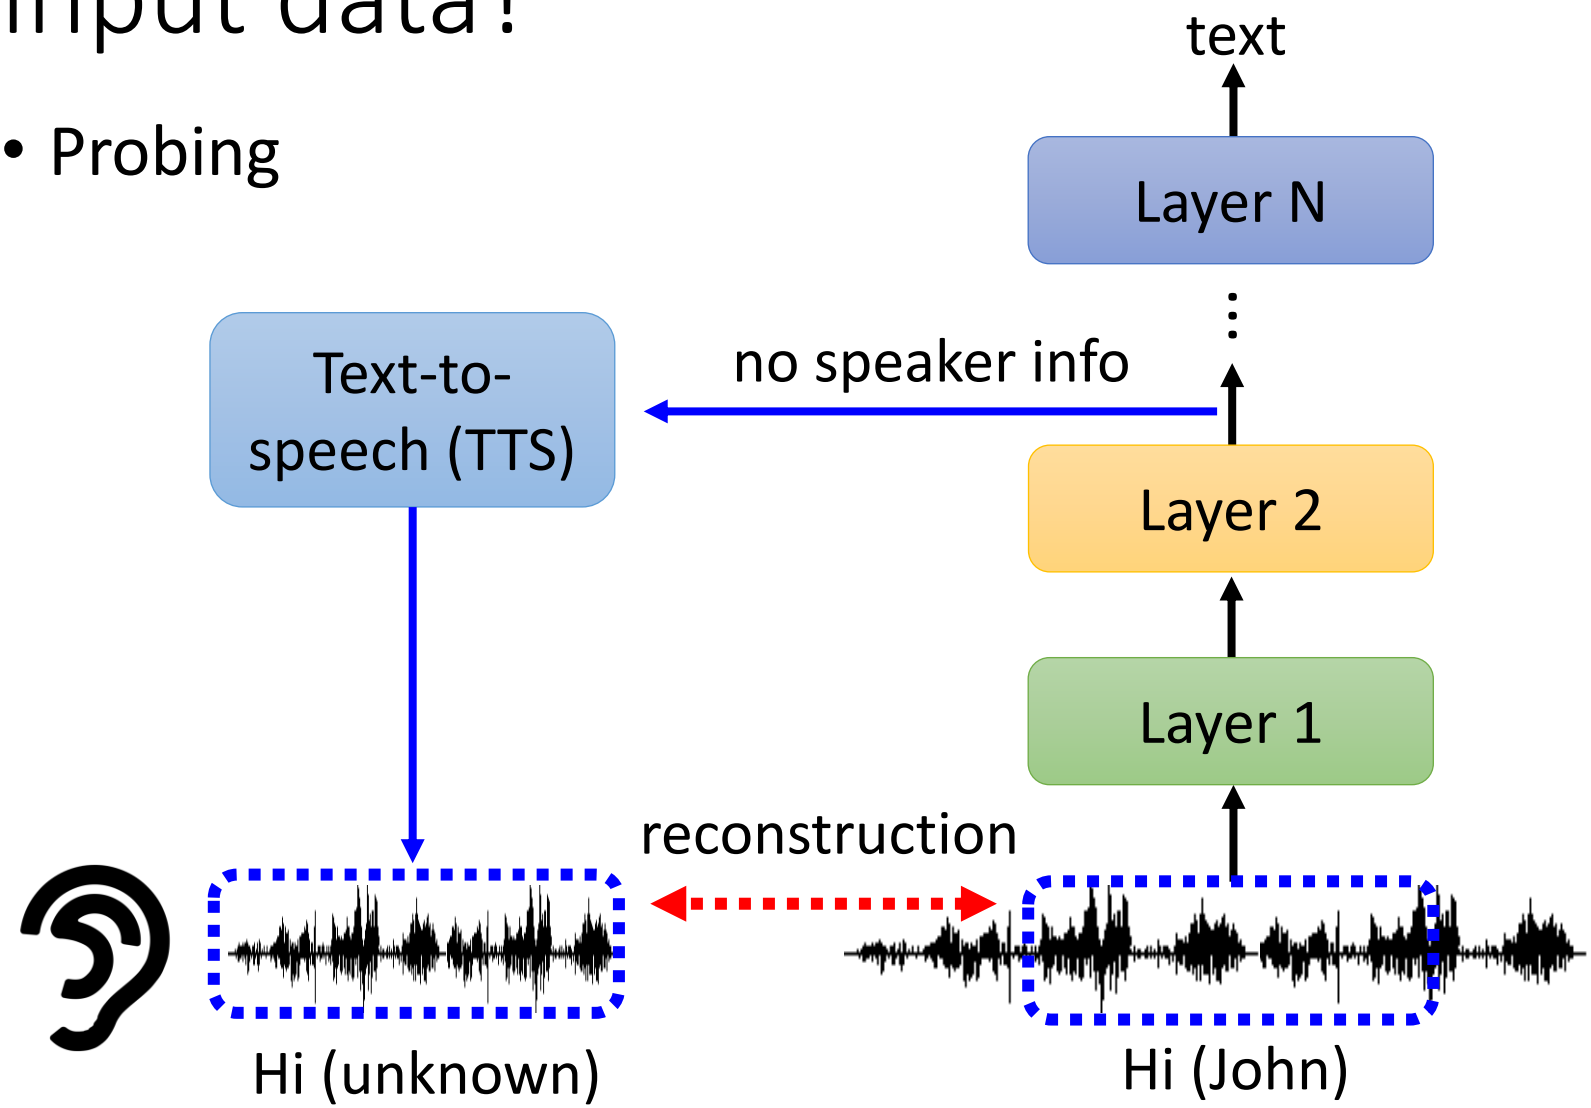

What does a network layer hear? Analyzing hidden representations of end-to-end ASR through speech synthesis

https://arxiv.org/abs/1911.01102

https://youtu.be/6gtn7H-pWr8

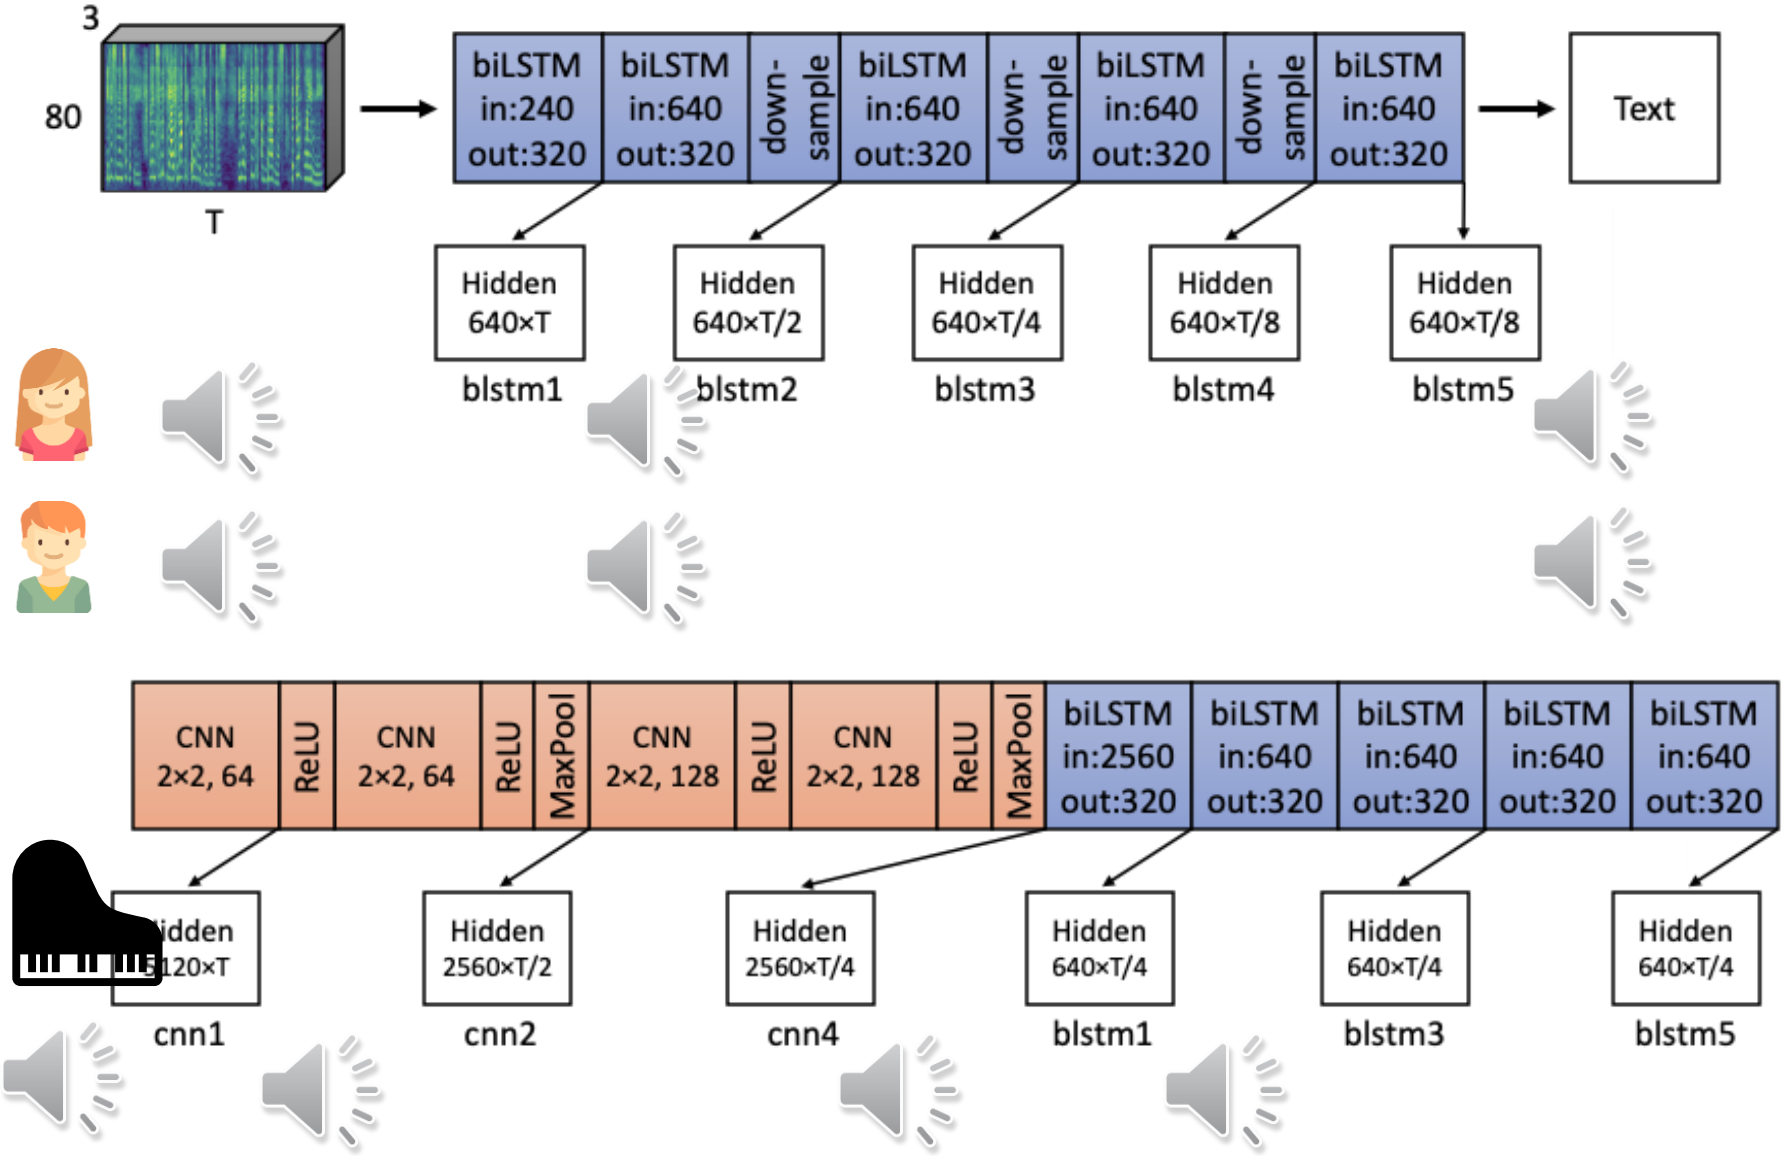

GLOBAL EXPLANATION: EXPLAIN THE WHOLE MODEL

Question: What does a “cat” look like?

What does a filter detect?

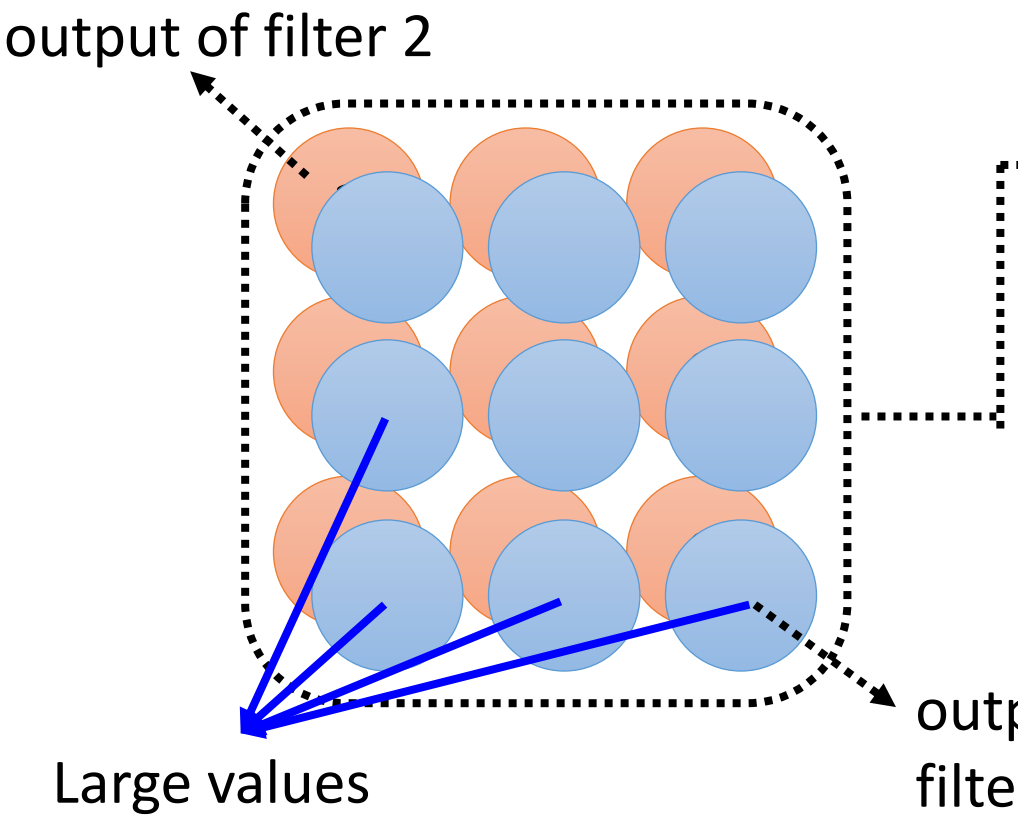

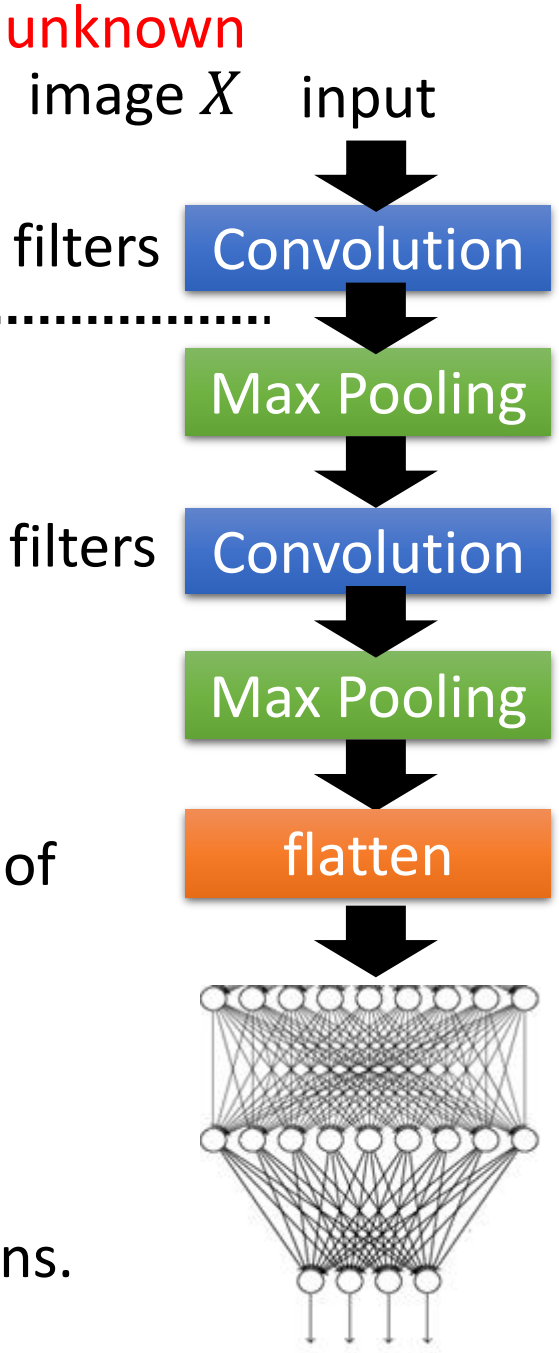

➡ Image X contains the patterns filter 1 can detect.

Let's **create** an image including the patterns.

What does a filter detect?

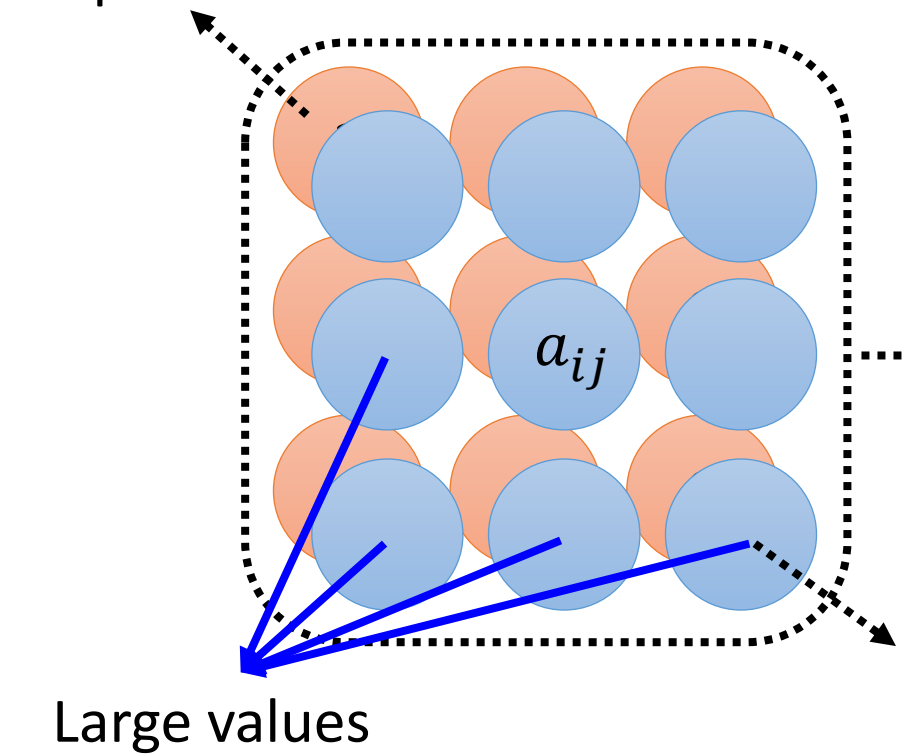

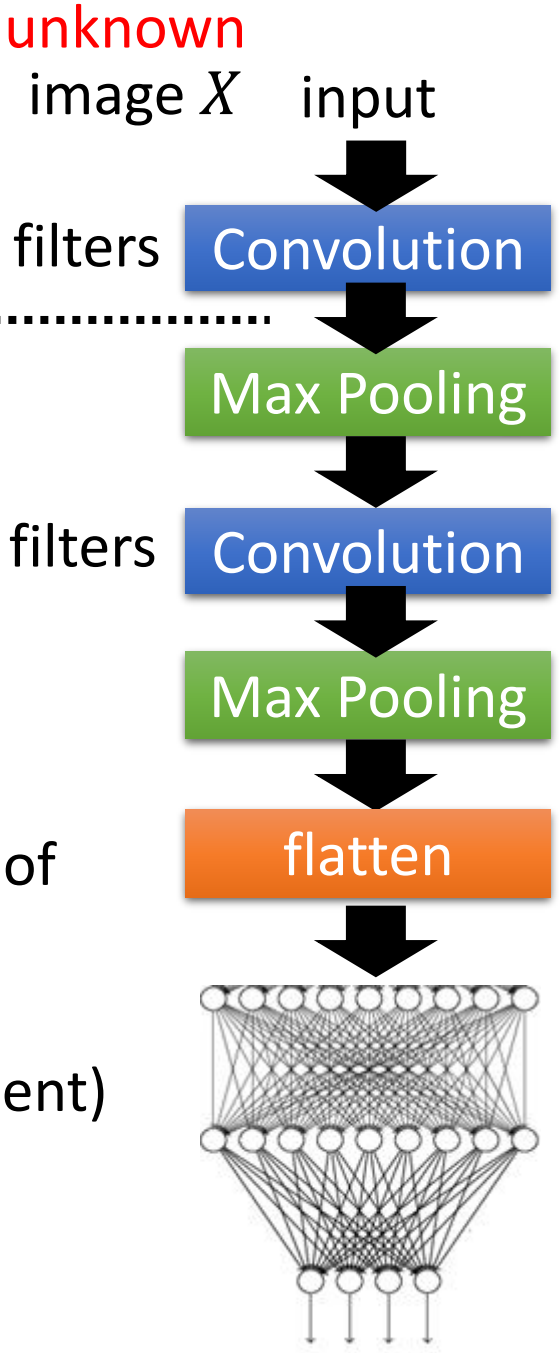

$$X^* = arg\max_X \sum_i \sum_j a_{ij}$$ (gradient ascent)

The image contains the patterns filter 1 can detect.

What does a filter detect?

E.g., Digit classifier

X^* for each filter

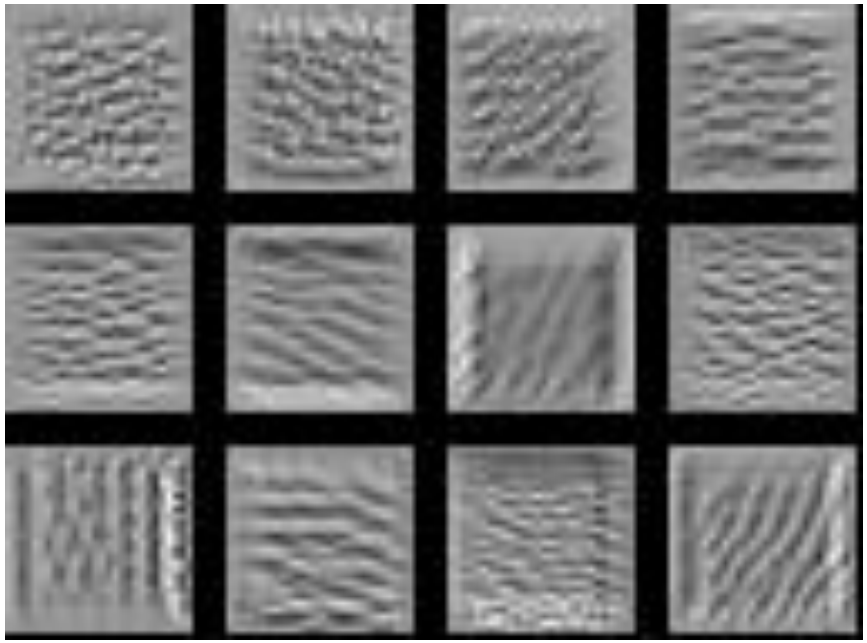

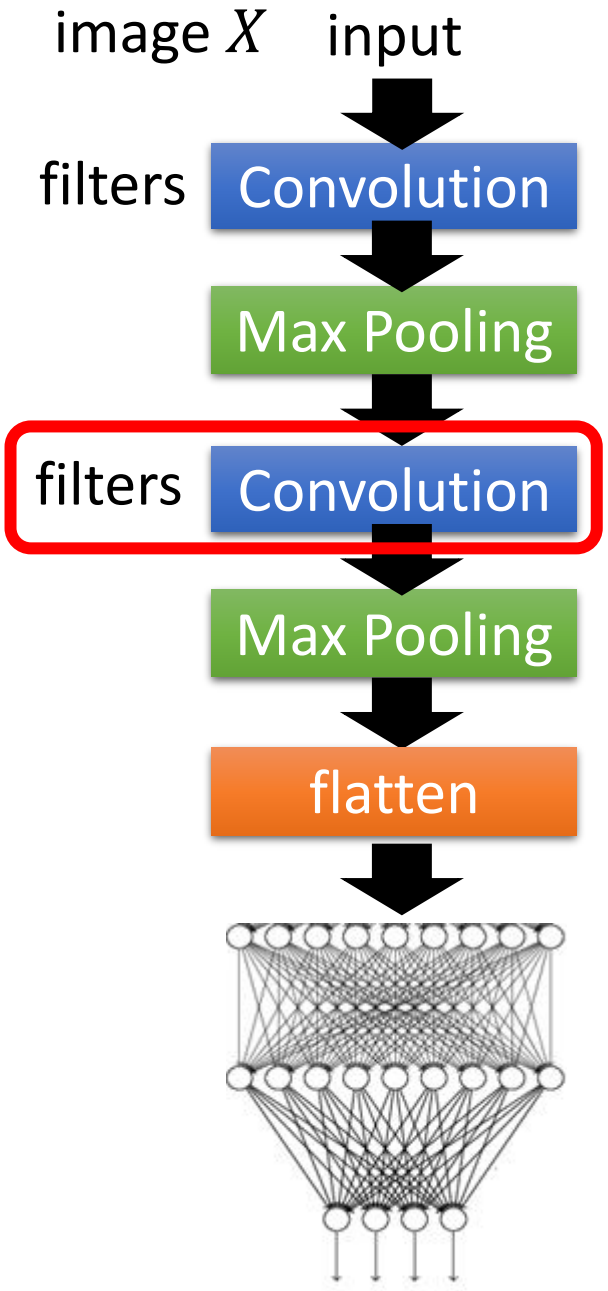

What does a digit look like for CNN?

E.g., Digit classifier

$$X^* = arg\max_X y_i$$ Can we see digits?

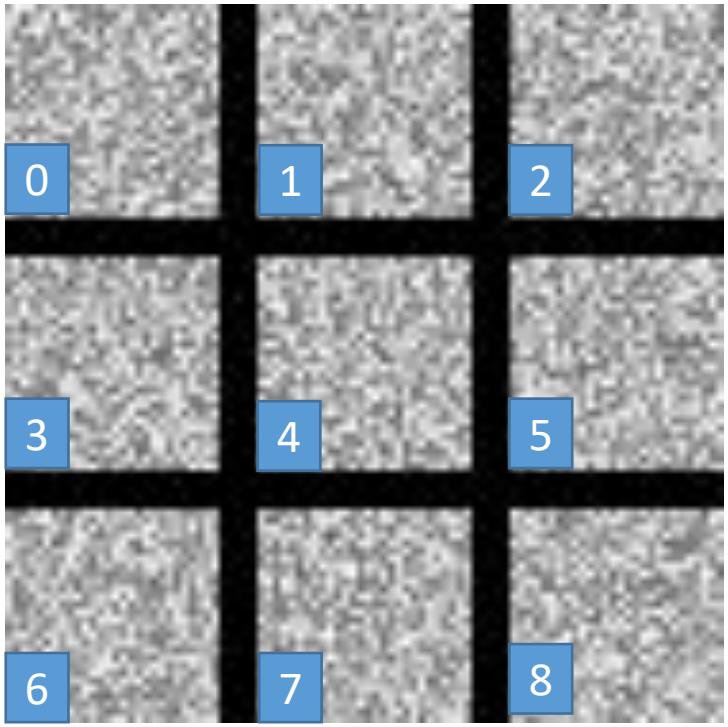

Surprise? Consider adversarial attack!

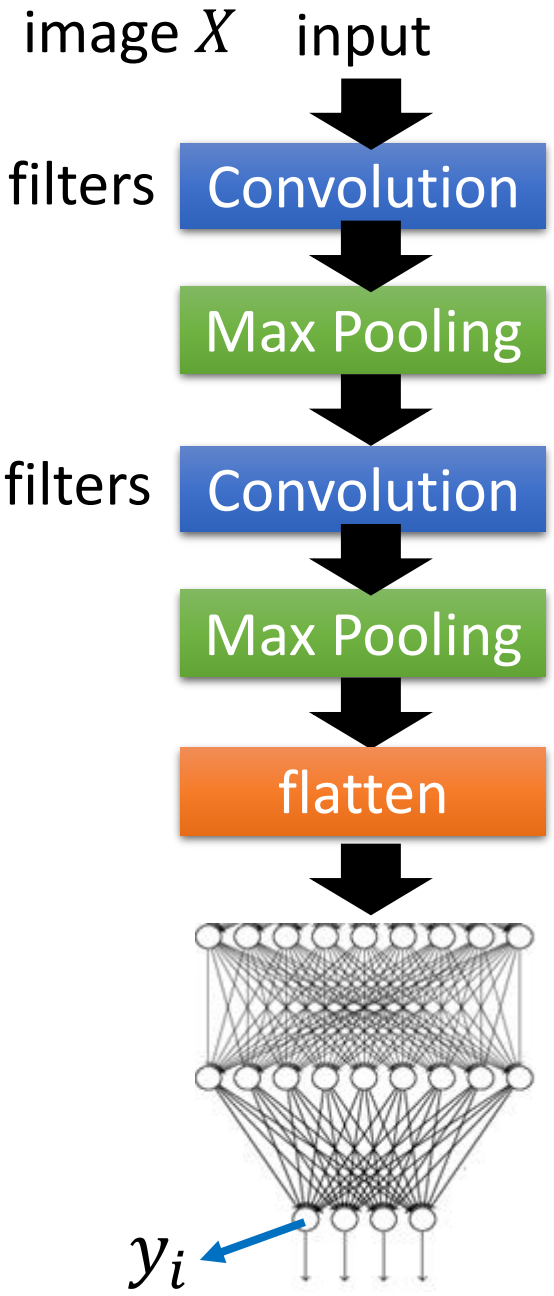

What does a digit look like for CNN?

Find the image that maximizes class probability

$$X^* = arg\max_X y_i$$

The image should looks like a digit.

$$X^* = arg\max_X y_i + R(X)$$

$$R(X) = -\sum_{i,j} |X_{ij}|$$

How likely X is a digit

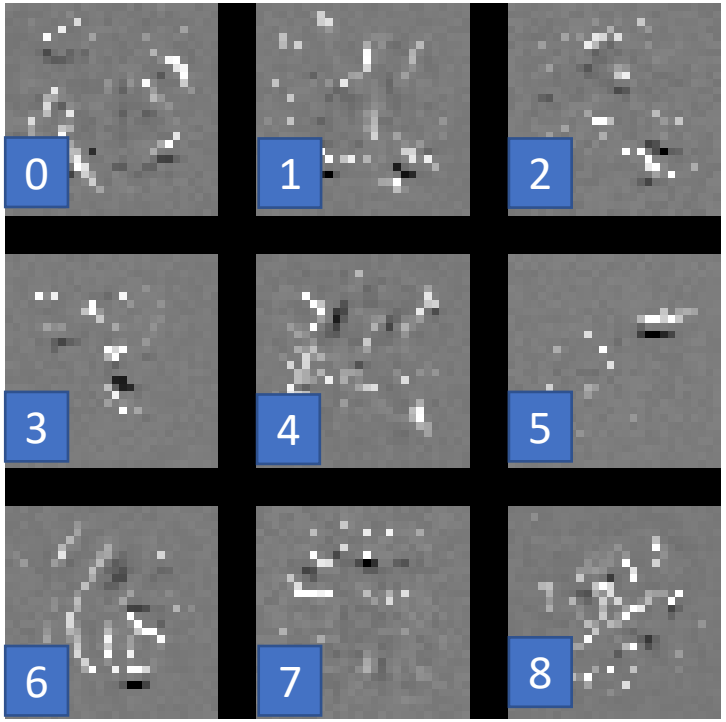

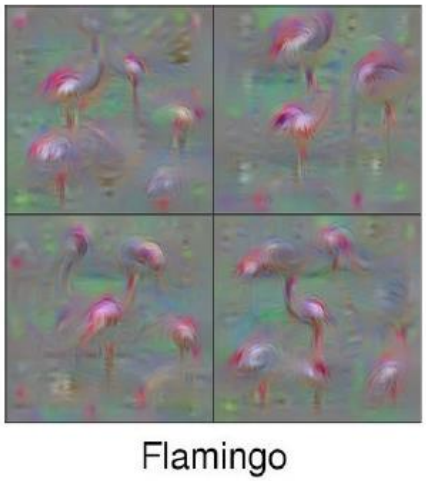
Flamingo

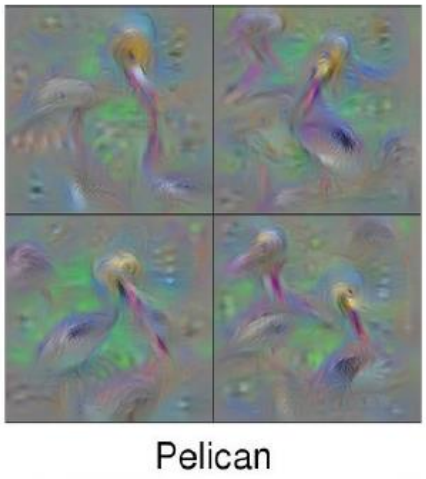
Pelican

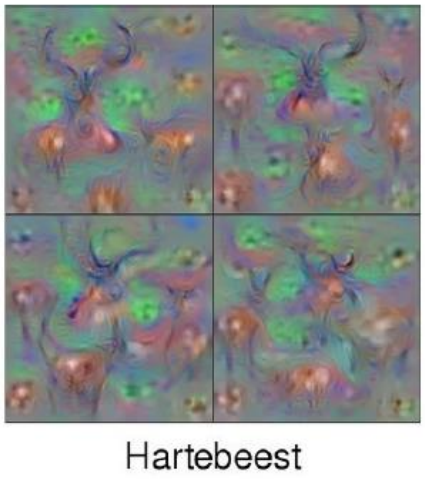
Hartebeest

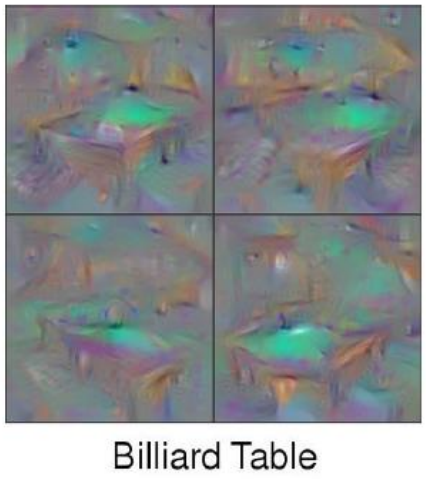
Billiard Table

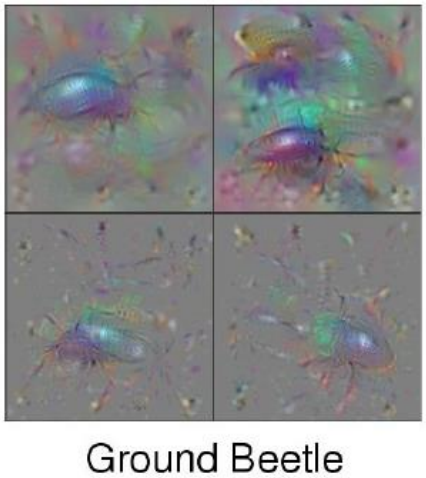
Ground Beetle

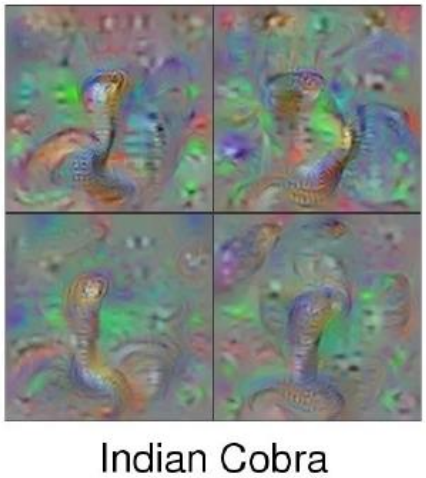
Indian Cobra

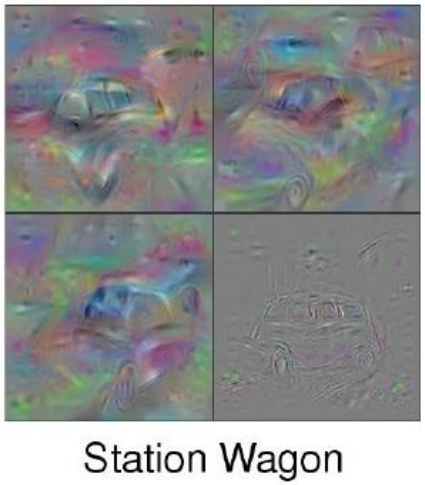
Station Wagon

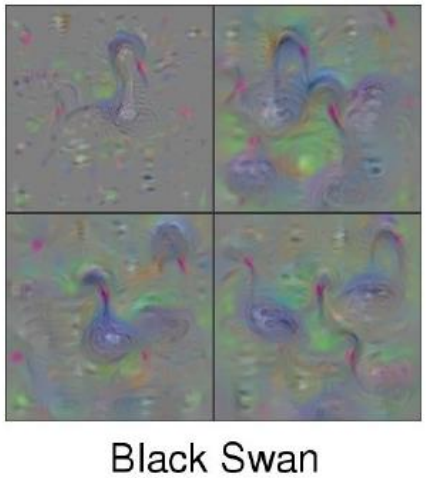
Black Swan

With several regularization terms, and hyperparameter tuning …..

https://arxiv.org/abs/1506.06579

(Simplified Version)

Constraint from Generator

- Training a generator (by GAN, VAE, etc.)

Training Examples

low-dim vector z → Image Generator G → Image X $X = G(z)$

z → Image Generator → Image X → Image Classifier → y

$$X^* = arg\max_X y_i \Rightarrow z^* = arg\max_z y_i$$

Show image: $X^* = G(z^*)$

https://arxiv.org/abs/1612.00005

Concluding Remarks

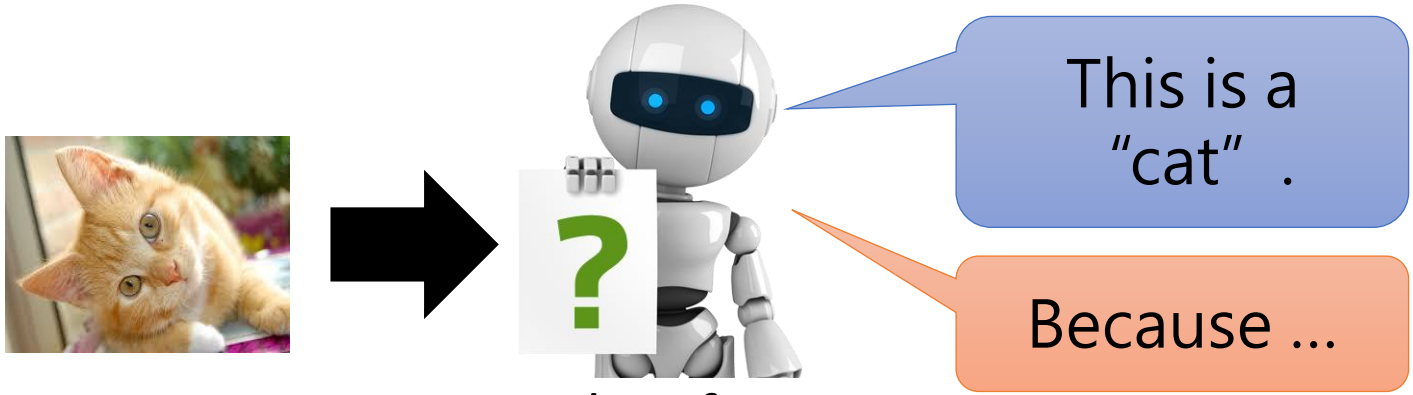

Local Explanation

Why do you think *this image* is a cat?

Global Explanation

What does a "cat" look like?

(not referred to a specific image)

Outlook

Using an interpretable model to mimic the behavior of an uninterpretable model.

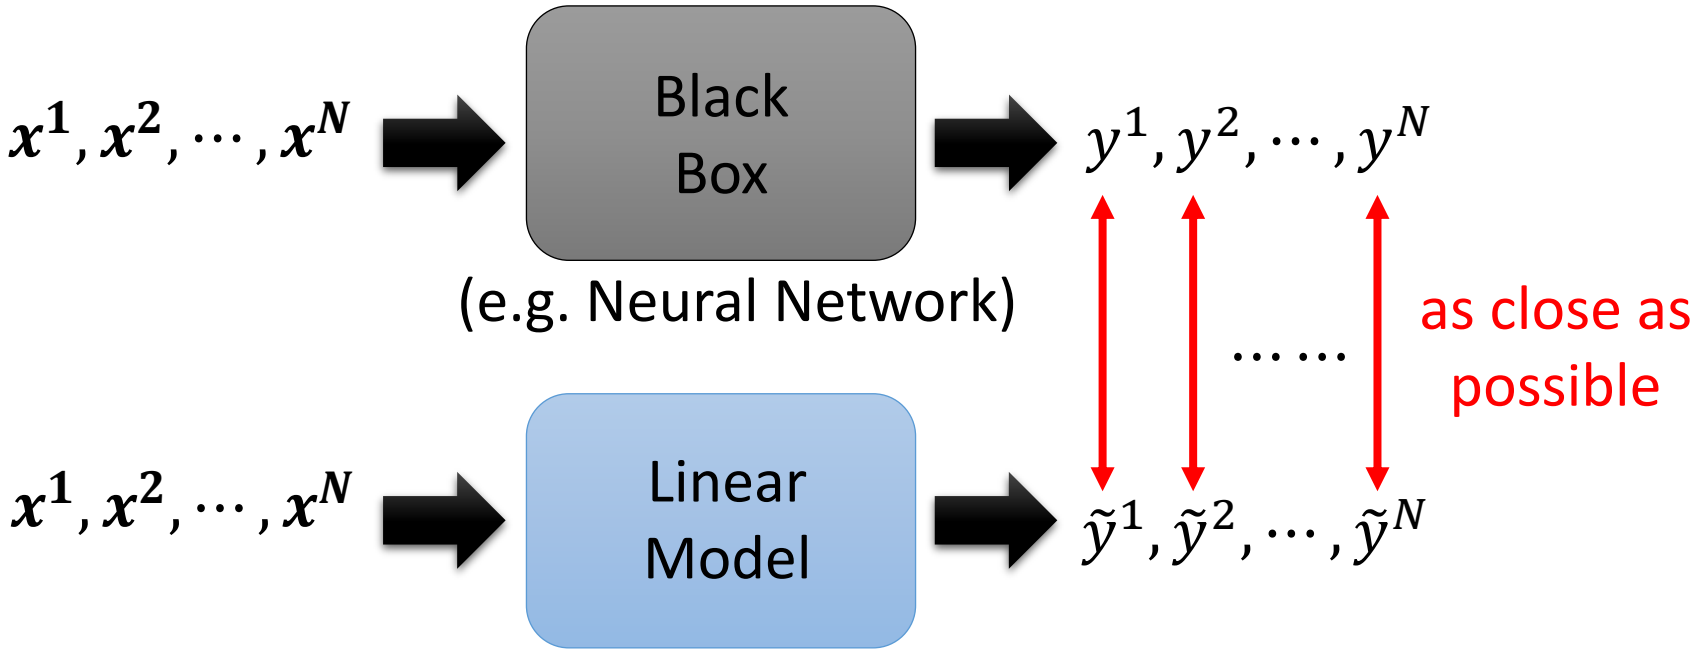

Local Interpretable Model-Agnostic Explanations (LIME)

https://youtu.be/K1mWgthGS-A
https://youtu.be/OjqIVSwly4k